

Home Prices EDA and Machine Learning

Executive Summary

As a team, we set out to develop a model to predict apartments prices in South Korea based on various features about the location and attributes of the particular apartment. Housing prices is an area of big concern all over the world given our ever growing population. South Korea, a country in which about 59.9% of all homes are apartments, is no exception. We will specifically be examining apartments in order to understand which factors contribute most to apartments prices, and to predict prices for apartments that we have not yet seen. For a family searching for a home in South Korea, this model can help them to understand if they are overpaying. For example, if they know the location of the apartment and some basic information, such as the number of rooms, they can determine if they are being offered a fair price. In addition, the landlord or seller of the property can use this model to calculate a fair price for the apartment.

Our Dataset contains the following features:

- transaction_real_price: the price that the apartment was sold at (target variable)
- key: increments by one based on the row number
- apartment_id: a unique identifier for the building in which the apartment is found
- transaction_year_month: the year and month the transaction took place
- transaction_date: the date on which the transaction took place
- year_of_completion: the year the apartment was built
- exclusive_use_area: the total floor area of the building
- floor: the floor number that the apartment building is located
- latitude: latitude of the apartment
- longitude: longitude of the apartment
- address_by_law: address represented numerically
- total_parking_capacity_in_site: the number of parking spots for the entire building complex in which the apartment is located (potentially there could be multiple separate buildings in the site)
- total_household_count_in_site: The number of separate households located in the apartment complex
- apartment_building_count_in_sites: the number of separate apartments building in the apartment complex
- tallest_building_in_sites: the number of floors of the tallest building in the apartment complex
- lowest_building_in_sites: the number of floors of the shortest building in the apartment complex
- heat_type: the type of heat available tot he apartment (individual, central, district)
- heat_fuel: the type of heating fuel used by the apartment (gas, cogeneration)
- room_id: unique identifier for the apartment
- supply_area: Total site area (area of the entire apartment complex)
- total_household_count_of_area_type: Count of households in the immediate area
- room_count: the number of rooms in the apartment
- bathroom_count: the number of bathrooms in the apartment
- front_door_structure: the structure of the entrance to the apartment (corridor, stairway, mixed)

EDA and Data Cleaning

Load in dataset

```
setwd("C:/Users/pmank/Dropbox/BU/BA810/group_project")
Price <- fread("trainPrice.csv")
```

First few rows

We can see that we have a total of 25 features in the dataset

```
head(Price, 5)
```

```
##      key apartment_id city transaction_year_month transaction_date
## 1:    0          5584   1              200601          11~20
## 2:    1          5584   1              200601          11~20
## 3:    2          5059   1              200601          11~20
## 4:    3          2816   1              200601          11~20
## 5:    4          2816   1              200601          11~20
##      year_of_completion exclusive_use_area floor latitude longitude
## 1:              1999              47.43    6 37.58597 127.0002
## 2:              1999              44.37    8 37.58597 127.0002
## 3:              1992              54.70    8 37.58051 127.0140
## 4:              1993              64.66   11 37.58032 127.0118
## 5:              1993             106.62    7 37.58032 127.0118
##      address_by_law total_parking_capacity_in_site total_household_count_in_sites
## 1:      1111017100              163              136
## 2:      1111017100              163              136
## 3:      1111017400              902              585
## 4:      1111017400              902              919
## 5:      1111017400              902              919
##      apartment_building_count_in_sites tallest_building_in_sites
## 1:              1              8
## 2:              1              8
## 3:              5             14
## 4:              7             15
## 5:              7             15
##      lowest_building_in_sites heat_type heat_fuel room_id supply_area
## 1:              4 individual      gas    91120      65.63
## 2:              4 individual      gas    91119      61.39
## 3:              9 individual      gas     8430      72.36
## 4:             11 individual      gas     5839      87.30
## 5:             11 individual      gas     5836     127.74
##      total_household_count_of_area_type room_count bathroom_count
## 1:              46              1              1
## 2:              10              2              1
## 3:             201              2              1
## 4:             284              2              1
## 5:             112              4              2
##      front_door_structure transaction_real_price
## 1:      corridor      215000000
## 2:      corridor      200000000
## 3:      corridor      168000000
## 4:      corridor      165000000
## 5:      stairway      280000000
```

Structure of the dataset

There are a total of 1,601,458 observations

```
str(Price)
```

```
## Classes 'data.table' and 'data.frame': 1601458 obs. of 25 variables:
## $ key : int 0 1 2 3 4 5 6 7 8 9 ...
## $ apartment_id : int 5584 5584 5059 2816 2816 2815 2815 9867 2818 2817
...
## $ city : int 1 1 1 1 1 1 1 1 1 1 ...
## $ transaction_year_month : int 200601 200601 200601 200601 200601 200601 200601
200601 200601 200601 ...
## $ transaction_date : chr "11~20" "11~20" "11~20" "11~20" ...
## $ year_of_completion : int 1999 1999 1992 1993 1993 2000 2000 2005 1999 2002
...
## $ exclusive_use_area : num 47.4 44.4 54.7 64.7 106.6 ...
## $ floor : int 6 8 8 11 7 9 13 10 18 12 ...
## $ latitude : num 37.6 37.6 37.6 37.6 37.6 ...
## $ longitude : num 127 127 127 127 127 ...
## $ address_by_law : integer64 1111017100 1111017100 1111017400 1111017400 1
111017400 1111018700 1111018700 1114016200 ...
## $ total_parking_capacity_in_site : num 163 163 902 902 902 ...
## $ total_household_count_in_sites : int 136 136 585 919 919 964 964 461 2282 5150 ...
## $ apartment_building_count_in_sites : int 1 1 5 7 7 12 12 9 19 42 ...
## $ tallest_building_in_sites : num 8 8 14 15 15 23 23 23 20 18 ...
## $ lowest_building_in_sites : num 4 4 9 11 11 10 10 6 8 11 ...
## $ heat_type : chr "individual" "individual" "individual" "individua
l" ...
## $ heat_fuel : chr "gas" "gas" "gas" "gas" ...
## $ room_id : int 91120 91119 8430 5839 5836 5831 5833 11862 5843 5
842 ...
## $ supply_area : num 65.6 61.4 72.4 87.3 127.7 ...
## $ total_household_count_of_area_type: int 46 10 201 284 112 454 207 82 576 864 ...
## $ room_count : num 1 2 2 2 4 3 3 3 3 3 ...
## $ bathroom_count : num 1 1 1 1 2 2 1 2 2 1 ...
## $ front_door_structure : chr "corridor" "corridor" "corridor" "corridor" ...
## $ transaction_real_price : integer64 215000000 200000000 168000000 165000000 28000
0000 415000000 267000000 415000000 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

```
transform(Price, transaction_real_price = as.numeric(transaction_real_price))
```

```

##          key apartment_id city transaction_year_month transaction_date
##      1:         0         5584      1              200601          11~20
##      2:         1         5584      1              200601          11~20
##      3:         2         5059      1              200601          11~20
##      4:         3         2816      1              200601          11~20
##      5:         4         2816      1              200601          11~20
##      ---
## 1601454: 1605344         11500      0              201810          21~31
## 1601455: 1605346         16686      1              201810          21~31
## 1601456: 1605356         22243      0              201810          21~31
## 1601457: 1605366          3686      1              201810          21~31
## 1601458: 1605373          2937      1              201810          21~31
##          year_of_completion exclusive_use_area floor latitude longitude
##      1:              1999              47.4300      6 37.58597 127.0002
##      2:              1999              44.3700      8 37.58597 127.0002
##      3:              1992              54.7000      8 37.58051 127.0140
##      4:              1993              64.6600     11 37.58032 127.0118
##      5:              1993             106.6200      7 37.58032 127.0118
##      ---
## 1601454:              1999             118.4700     14 35.15557 129.0175
## 1601455:              2007              59.9900      4 37.50239 126.9420
## 1601456:              2014              84.9669     31 35.06480 128.9831
## 1601457:              1996              59.3400      4 37.55521 127.1313
## 1601458:              1999              84.8800      5 37.60433 127.0172
##          address_by_law total_parking_capacity_in_site
##      1:      1111017100              163
##      2:      1111017100              163
##      3:      1111017400              902
##      4:      1111017400              902
##      5:      1111017400              902
##      ---
## 1601454:      2623011100              876
## 1601455:      1159010200             1651
## 1601456:      2638010600             1761
## 1601457:      1174010700              111
## 1601458:      1129013300              802
##          total_household_count_in_sites apartment_building_count_in_sites
##      1:              136              1
##      2:              136              1
##      3:              585              5
##      4:              919              7
##      5:              919              7
##      ---
## 1601454:              819              8
## 1601455:             1122             22
## 1601456:             1326              9
## 1601457:              107              1
## 1601458:             860              8
##          tallest_building_in_sites lowest_building_in_sites  heat_type
##      1:              8              4 individual
##      2:              8              4 individual
##      3:             14              9 individual
##      4:             15             11 individual

```

```
##      5:      15      11 individual
##      ---
## 1601454:      27      13 individual
## 1601455:      15       8 individual
## 1601456:      35      27 individual
## 1601457:      19      11 individual
## 1601458:      22       7 individual
##      heat_fuel room_id supply_area total_household_count_of_area_type
##      1:      gas   91120      65.63              46
##      2:      gas   91119      61.39              10
##      3:      gas    8430      72.36             201
##      4:      gas    5839      87.30             284
##      5:      gas    5836     127.74             112
##      ---
## 1601454:      gas   44386     143.45             108
## 1601455:      gas   13884      79.98             254
## 1601456:      gas   56043     109.77             209
## 1601457:      gas  165820      88.37              4
## 1601458:      gas    6279     108.75             209
##      room_count bathroom_count front_door_structure transaction_real_price
##      1:         1              1      corridor      2.15e+08
##      2:         2              1      corridor      2.00e+08
##      3:         2              1      corridor      1.68e+08
##      4:         2              1      corridor      1.65e+08
##      5:         4              2      stairway      2.80e+08
##      ---
## 1601454:         4              2      stairway      4.27e+08
## 1601455:         3              2      stairway      7.71e+08
## 1601456:         3              2      stairway      3.43e+08
## 1601457:         3              1      corridor      4.85e+08
## 1601458:         3              2      stairway      4.30e+08
```

```
#summary(Price)
```

Data Cleaning

View the count of null values in each column

The feature, total_parking_capacity_in_site has the largest number of null values (91813)

```
Price[, lapply(.SD, function(x) sum(is.na(x)))]
```

```
##      key apartment_id city transaction_year_month transaction_date
## 1:      0              0      0                      0              0
##      year_of_completion exclusive_use_area floor latitude longitude
## 1:              0              0      0              0              0
##      address_by_law total_parking_capacity_in_site total_household_count_in_sites
## 1:              0                      91813                      0
##      apartment_building_count_in_sites tallest_building_in_sites
## 1:              0                      9
##      lowest_building_in_sites heat_type heat_fuel room_id supply_area
## 1:              9              0              0              0              0
##      total_household_count_of_area_type room_count bathroom_count
## 1:              0              691              691
##      front_door_structure transaction_real_price
## 1:              0              0
```

View all unique values present in columns

```
unique(Price$heat_type)
```

```
## [1] "individual" "central"    "district"   ""
```

How many values are empty strings or dashes?

First, we will view the unique values to understand if there are any other invalid values other than NA

```
unique(Price$city)
```

```
## [1] 1 0
```

```
unique(Price$bathroom_count)
```

```
## [1] 1 2 0 NA 3 4 5
```

```
unique(Price$room_count)
```

```
## [1] 1 2 4 3 5 0 6 NA 8 7
```

```
unique(Price$front_door_structure)
```

```
## [1] "corridor" "stairway" "mixed"     ""          "-"
```

```
unique(Price$year_of_completion)
```

```
## [1] 1999 1992 1993 2000 2005 2002 2001 1997 1996 1990 1989 1987 1985 1995 1988
## [16] 1991 1977 1971 1998 1974 1994 2003 1984 1986 1982 1983 2004 1975 1981 1980
## [31] 1978 1976 1979 2006 1973 1962 1970 1968 1969 1972 2007 2008 2009 1966 2010
## [46] 2011 2012 2013 2014 2015 2016 2017 2018
```

A number of columns have an empty string value or one dash rather than an NA value, we will remove these from the dataset.

```
Price[heat_type == ''][, .N]
```

```
## [1] 2017
```

```
Price[heat_fuel == ''][, .N]
```

```
## [1] 9667
```

```
Price[front_door_structure == ''][, .N]
```

```
## [1] 13892
```

```
Price[heat_fuel == '-'][, .N]
```

```
## [1] 8971
```

```
Price[front_door_structure == '-'][, .N]
```

```
## [1] 21
```

We can remove these empty string and dash values

```
Price <- Price[heat_fuel != '']
Price <- Price[heat_fuel != '-']
Price <- Price[front_door_structure != '-']
Price <- Price[heat_type != '']
Price <- Price[front_door_structure != '']
```

Drop rows room_count as 8

After examining the dataset, we determined that this was an outlier since only a very small number of apartments were shown as having 8 rooms

```
Price <- Price[room_count != 8]
```

Remove all rows that have missing values

There are a total of 1601458 million observations, and we will be removing less than 100,000 of them. We can confirm that rows were removed using .N (the number of observations)

```
Price <- na.omit(Price)
Price[, .N]
```

```
## [1] 1478931
```

Create dummy variables for use in later ML steps

We will convert transaction_date, heat_type, heat_fuel, and front_door_structure

```
Price <- Price[transaction_date == "1~10", transaction_date:=1]
Price <- Price[transaction_date == "11~20", transaction_date:=2]
Price <- Price[transaction_date == "21~30", transaction_date:=3]
Price <- Price[transaction_date == "21~28", transaction_date:=3]
Price <- Price[transaction_date == "21~29", transaction_date:=3]
Price <- Price[transaction_date == "21~31", transaction_date:=3]
Price[,transaction_date := as.numeric(transaction_date)]

Price <- Price[heat_fuel == "gas", heat_fuel:=0]
Price <- Price[heat_fuel == "cogeneration", heat_fuel:=1]
Price[,heat_fuel := as.numeric(heat_fuel)]

Price[,transaction_real_price := as.numeric(transaction_real_price)]
Price[,address_by_law := as.numeric(address_by_law)]

Price <- fastDummies::dummy_cols(Price)
```

EDA

In the EDA section, we wanted to create a few graphs to help us understand the distribution of our features, and also how our features relate to the target variable, transaction_real_price (this is the price that each apartments sells at). The charts below reveal that there are a number of features that seem to relate to the price, including the city, number of rooms, front door structure, heating fuel type, and heat type. These charts were also helpful in understanding unusual values. We found that there were a few apartments that had 8 rooms, but a very low price. We decided to remove these above in the data cleaning section.

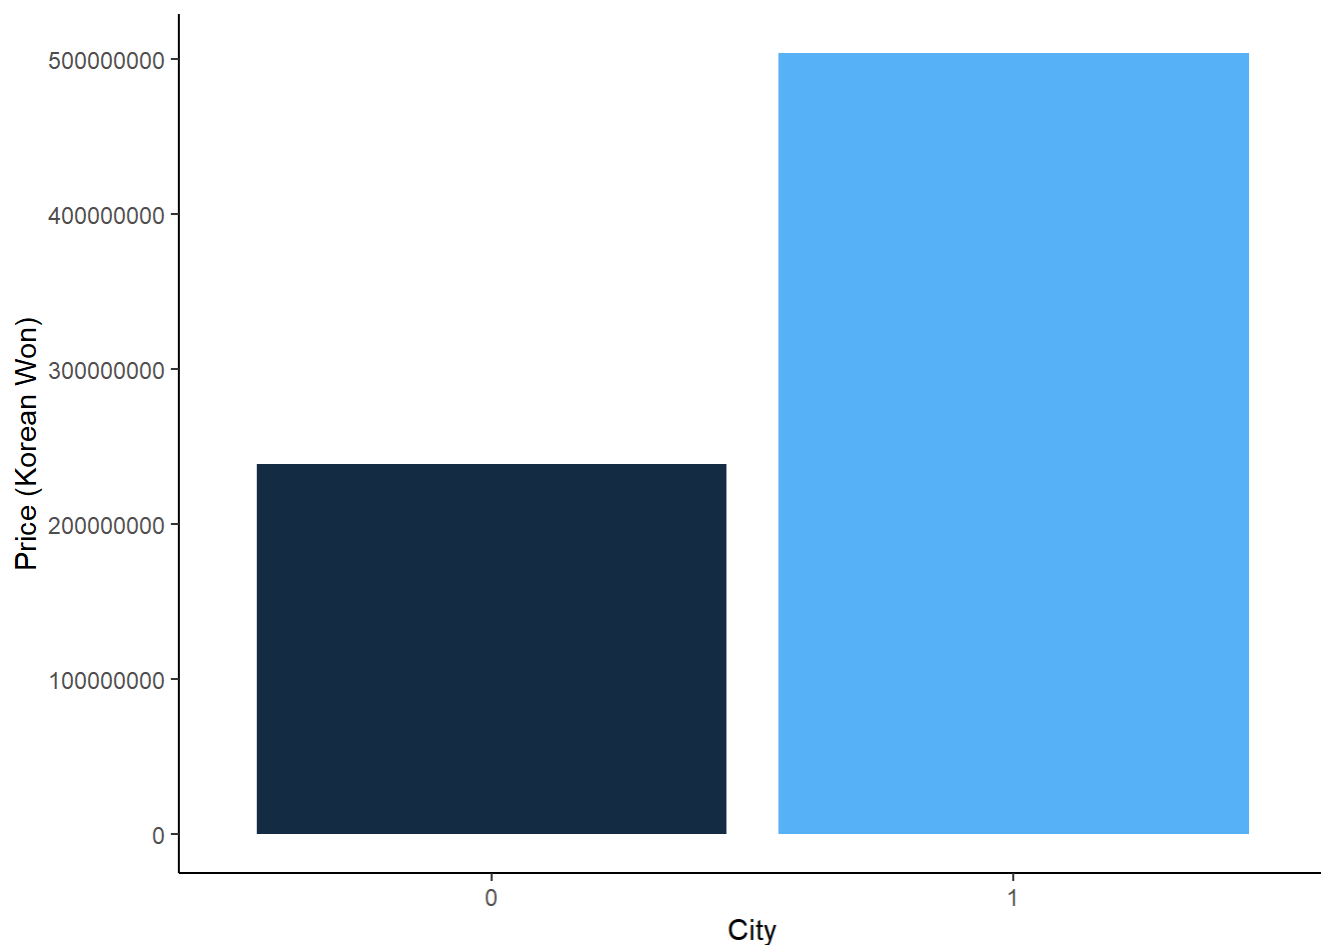
Barchart showing real price and city

Below, 1 is Busan and 0 is Seoul. We can see that on average, prices tend to be higher in Busan.

```
ggplot(Price, aes(x=as.factor(city), y = transaction_real_price, fill = city)) +

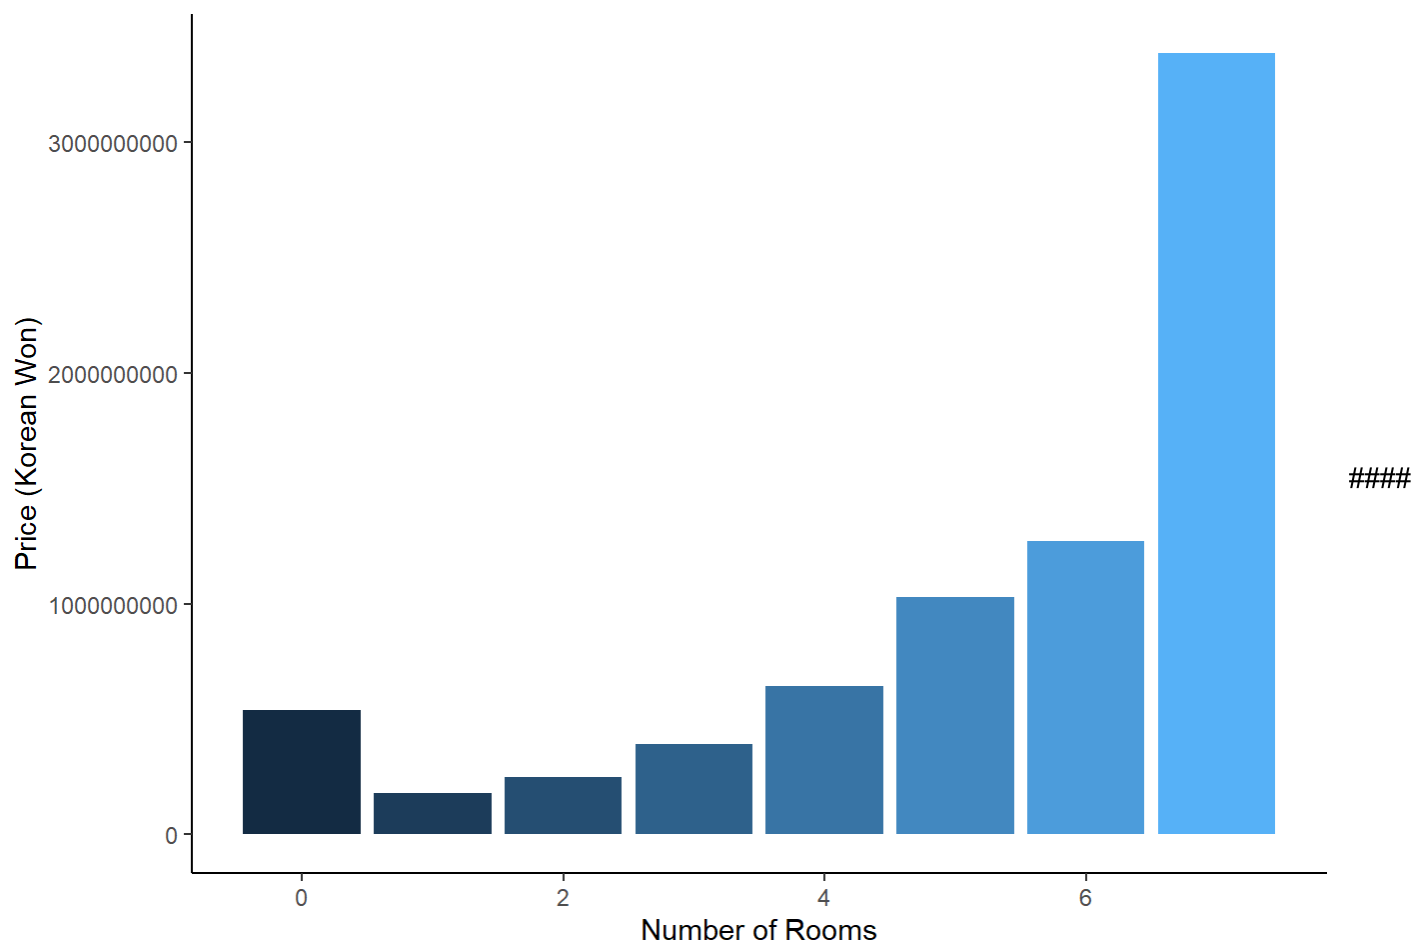
  geom_bar(stat = "summary", fun = "mean") +
  scale_color_manual(values=c("red", "blue")) +
  scale_y_continuous(labels = function(x) format(x, scientific = FALSE)) +

  theme(legend.position="none") +
  xlab('City') +
  ylab('Price (Korean Won)')
```

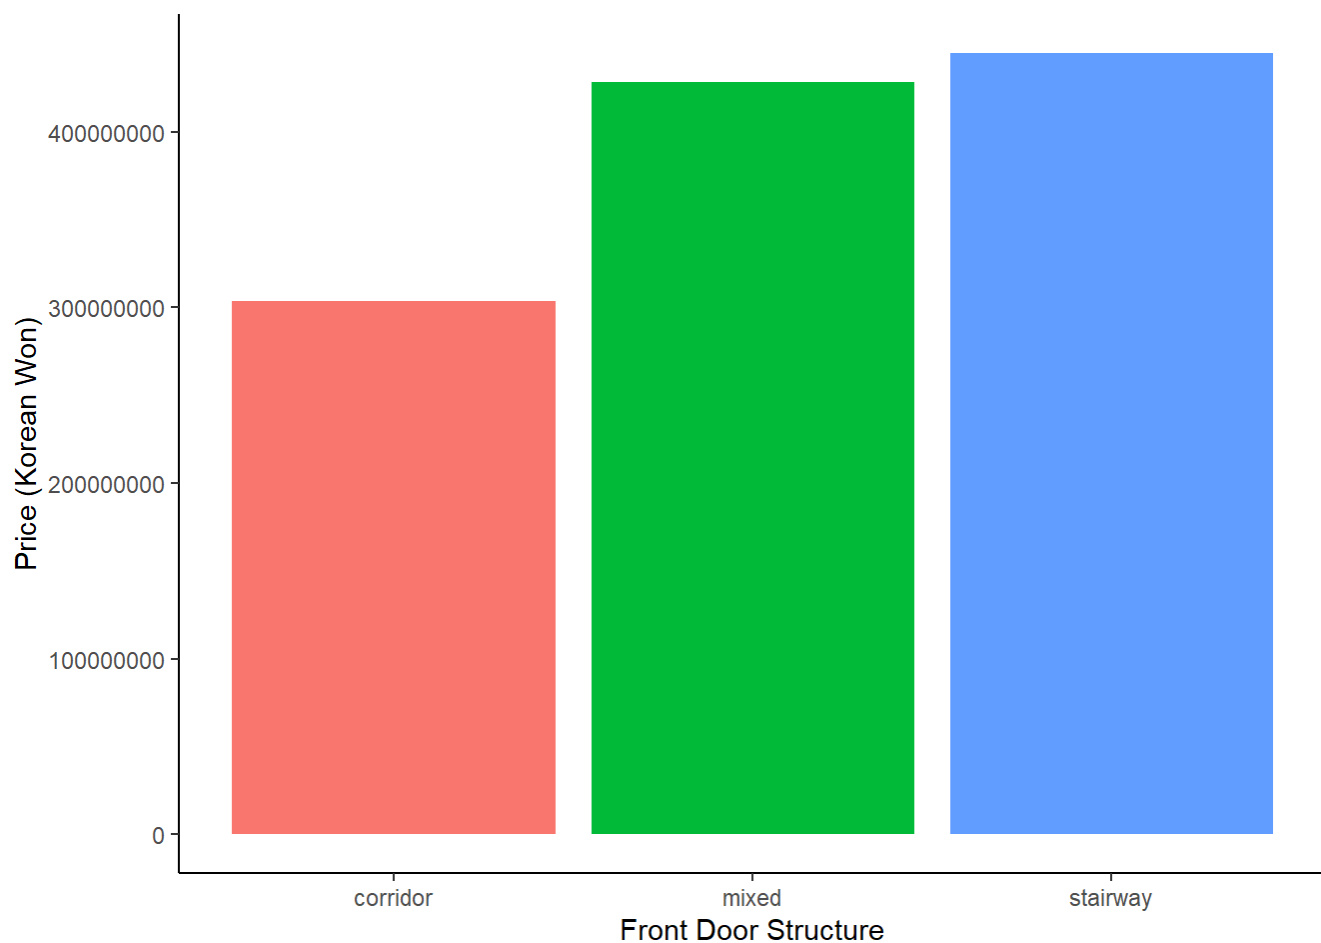
Bar chart showing real price and room count

```
ggplot(Price, aes(x=room_count, y = transaction_real_price, fill = room_count)) +  
  geom_bar(stat = "summary", fun = "mean")+  
  scale_y_continuous(labels = function(x) format(x, scientific = FALSE)) +  
  theme(legend.position="none")+  
  xlab('Number of Rooms') +  
  ylab('Price (Korean Won)')
```



Bar chart showing real price and front_door_structure

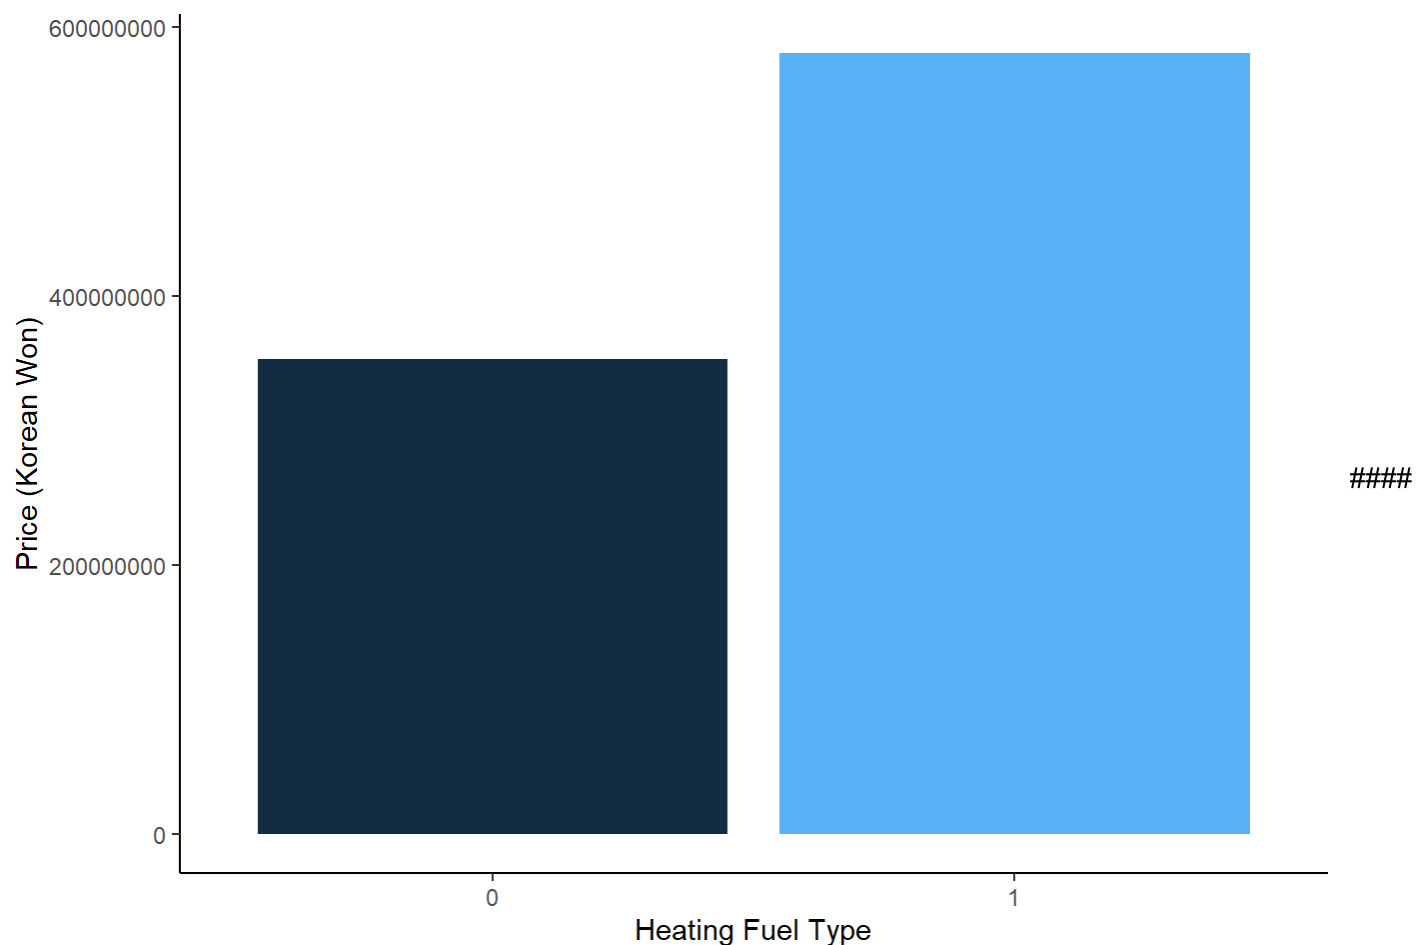
```
ggplot(Price, aes(x=front_door_structure, y = transaction_real_price, fill = front_door_structur
e)) +
  geom_bar(stat = "summary", fun = "mean")+
  scale_y_continuous(labels = function(x) format(x, scientific = FALSE)) +
  theme(legend.position="none")+
  xlab('Front Door Structure') +
  ylab('Price (Korean Won)')
```



Bar chart showing real price and heat_fuel

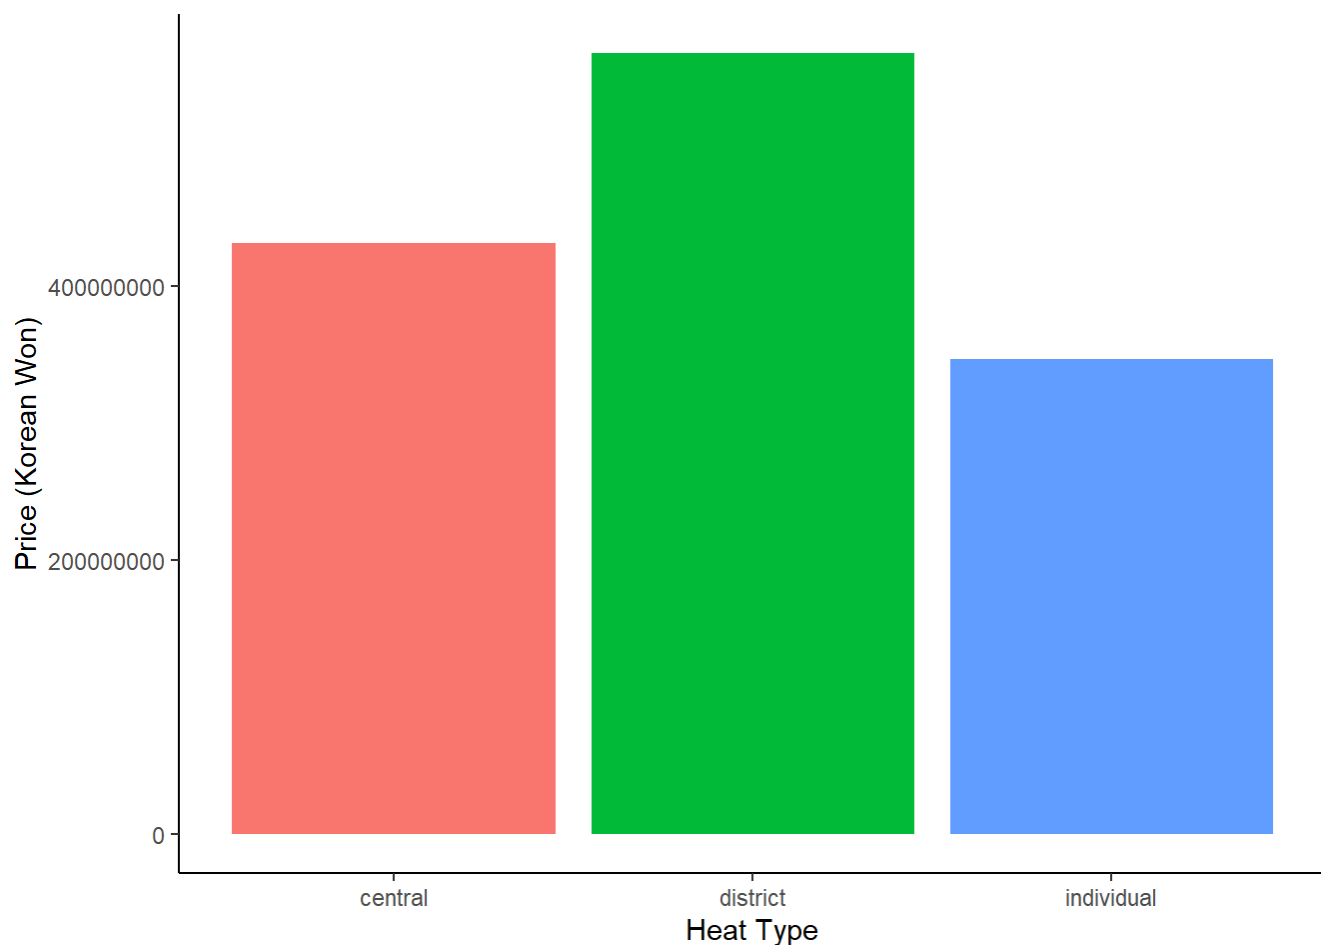
Below, 0 represents gas and 1 is co-generation

```
ggplot(Price, aes(x=as.factor(heat_fuel), y = transaction_real_price, fill = heat_fuel)) +  
  geom_bar(stat = "summary", fun = "mean")+  
  scale_y_continuous(labels = function(x) format(x, scientific = FALSE)) +  
  theme(legend.position="none")+  
  xlab('Heating Fuel Type') +  
  ylab('Price (Korean Won)')
```



Bar chart showing real price and heat_type Below, 0 is individual, 1 is central and 2 is district

```
ggplot(Price, aes(x=heat_type, y = transaction_real_price, fill = heat_type)) +  
  geom_bar(stat = "summary", fun = "mean")+  
  scale_y_continuous(labels = function(x) format(x, scientific = FALSE)) +  
  theme(legend.position="none")+  
  xlab('Heat Type') +  
  ylab('Price (Korean Won)')
```



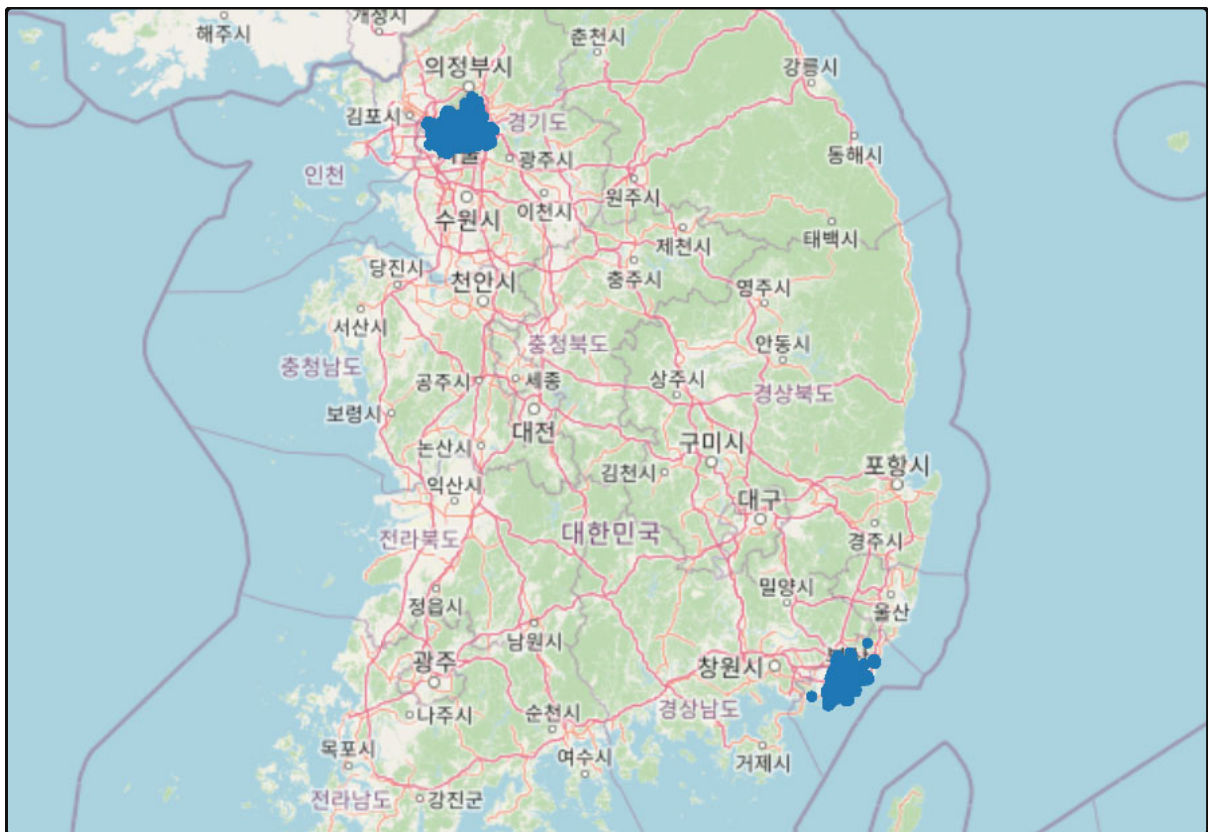
Plotly Map showing locations of apartments

We can see that all of the apartments are either in Seoul or Busan. This is coded in the data in the city feature.

```
fig <- head(Price, 100000)
fig <- fig %>%
  plot_ly(
    lat = ~latitude,
    lon = ~longitude,
    #marker = list(color = "fuchsia"),
    color = Price["transaction_real_price"],
    type = 'scattermapbox',
    hovertext = Price["transaction_real_price"])
fig <- fig %>%
  layout(
    mapbox = list(
      style = 'open-street-map',
      zoom = 2.5,
      center = list(longitude = 35.963911, latitude = 127.919770)))

fig
```

```
## No scattermapbox mode specified:
##   Setting the mode to markers
##   Read more about this attribute -> https://plotly.com/r/reference/#scatter-mode
```



Machine Learning

In this section, we used a number of machine learning techniques to predict the prices of apartments in South Korea. We used the following models:

- Linear Regression
- Forward Regression
- Backward Regression
- Ridge Regression
- Lasso Regression
- Decision Tree
- Bagging
- Random Forest
- Boosting

Drop useless columns

Below, we will remove `apartment_id`, `room_id`, and `key` since these are not useful in prediction. We also removed redundant dummy variables.

```
Price[,key := NULL]
Price[,apartment_id := NULL]
Price[,room_id := NULL]
Price[,heat_type := NULL]
Price[,front_door_structure := NULL]
Price[,heat_type_individual := NULL]
Price[,front_door_structure_stairway := NULL]
# prevent aliased coefficients
```

Setup test and train datasets

```
# Set the seed to get consistent results
set.seed(810)

rows <- sample(nrow(Price), 160000)
Price <- Price[rows,]

# Split the data
# This way we can do an 80/20 split
row_index <- sample(nrow(Price), 128000)

# we use that set of random numbers to select those random rows
dd_train <- Price[row_index,]
dd_test <- Price[-row_index,]
```

Linear Regression

We started by trying linear regression model. We used an 80/20 split between our test and train datasets, and used all of the features to predict price. For this model, we calculated a Train RMSE score of 188,332,311 Korean Won. We calculated a Test RMSE score of 188,544,540 Korean Won.

```
# our response variables to use later
set.seed(810)
y_train <- dd_train$transaction_real_price
y_test <- dd_test$transaction_real_price

# fit the full model
fit_lm1 <- lm(transaction_real_price ~ ., data=dd_train)
yhat_train_lm1 <- predict(fit_lm1)
mse_train_lm1 <- mean((y_train - yhat_train_lm1)^2)
paste("Linear Regression Train RMSE",sqrt(mse_train_lm1))
```

```
## [1] "Linear Regression Train RMSE 188332311.411548"
```

```
yhat_test_lm1 <- predict(fit_lm1, dd_test)
mse_test_lm1 <- mean((y_test - yhat_test_lm1)**2)
paste("Linear Regression Test RMSE",sqrt(mse_test_lm1))
```

```
## [1] "Linear Regression Test RMSE 188544540.361935"
```

A summary of the coefficient values in the model.

```
summary(fit_lm1)
```

```
##
## Call:
## lm(formula = transaction_real_price ~ ., data = dd_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.118e+09 -9.743e+07 -1.692e+07  7.074e+07  6.097e+09
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -5.436e+10  9.421e+08  -57.707   < 2e-16 ***
## city           5.765e+09  5.642e+07  102.180   < 2e-16 ***
## transaction_year_month 1.668e+05  1.428e+03  116.823   < 2e-16 ***
## transaction_date    6.465e+05  6.468e+05   1.000 0.317529
## year_of_completion  -2.637e+06  9.746e+04  -27.053   < 2e-16 ***
## exclusive_use_area   1.553e+06  1.346e+05  11.538   < 2e-16 ***
## floor           1.869e+06  8.684e+04  21.525   < 2e-16 ***
## latitude        -1.026e+09  1.106e+07  -92.799   < 2e-16 ***
## longitude         4.498e+08  7.047e+06  63.830   < 2e-16 ***
## address_by_law      1.417e+00  3.967e-02  35.713   < 2e-16 ***
## total_parking_capacity_in_site 2.850e+04  1.047e+03  27.226   < 2e-16 ***
## total_household_count_in_sites -8.174e+04  1.679e+03  -48.684   < 2e-16 ***
## apartment_building_count_in_sites 6.668e+06  8.495e+04  78.494   < 2e-16 ***
## tallest_building_in_sites 2.406e+06  1.213e+05  19.832   < 2e-16 ***
## lowest_building_in_sites 3.271e+06  1.199e+05  27.291   < 2e-16 ***
## heat_fuel         5.167e+07  3.737e+06  13.828   < 2e-16 ***
## supply_area        3.326e+06  1.148e+05  28.977   < 2e-16 ***
## total_household_count_of_area_type -1.459e+04  1.929e+03  -7.565 3.90e-14 ***
## room_count        -6.117e+06  1.280e+06  -4.778 1.77e-06 ***
## bathroom_count     -6.561e+06  1.725e+06  -3.803 0.000143 ***
## heat_type_central    7.847e+06  2.117e+06   3.707 0.000210 ***
## heat_type_district   6.760e+07  3.779e+06  17.890   < 2e-16 ***
## front_door_structure_corridor -1.441e+07  1.760e+06  -8.184 2.77e-16 ***
## front_door_structure_mixed -2.789e+06  4.178e+06  -0.668 0.504395
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 188300000 on 127976 degrees of freedom
## Multiple R-squared:  0.6566, Adjusted R-squared:  0.6565
## F-statistic: 1.064e+04 on 23 and 127976 DF, p-value: < 2.2e-16
```

Forward Selection

In forward selection, we used cross validation to split our data into 10 groups. In this model, we found that as the number of predictors in the model increased, our MSE decreased. This model would suggest that all of the features would be useful; however, we know that forward and backward selection don't test every possible combination.

In the output below, we can see that there is a separate results for each model, starting with just one explanatory variable, and ending with a model with all 24 variables.

In this model, our lowest RMSE was 188,321,979 Korean Won with 21 variables.

```
# Set up repeated k-fold cross-validation
set.seed(810)
train.control <- trainControl(method = "cv", number = 10)

# Train the model

step.model_f <- train(transaction_real_price ~., data = dd_train,
  method = "leapForward",
  tuneGrid = data.frame(nvmax = 1:23),
  trControl = train.control
)
step.model_f$results
```

##	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	275820794	0.2633316	183550247	5271468	0.012969232	1810020.1
## 2	2	234909141	0.4656794	153777481	5215660	0.009970077	1767563.1
## 3	3	223937947	0.5144075	145692662	4701072	0.007960067	1515730.5
## 4	4	215028464	0.5523260	137104337	4926310	0.007136277	1547467.7
## 5	5	206779570	0.5860626	129796361	4932603	0.006514049	1288489.4
## 6	6	200154326	0.6121798	126424032	5129113	0.007657021	1145261.7
## 7	7	194883374	0.6323292	123613028	5326193	0.008018565	1179427.1
## 8	8	193296484	0.6382953	122648186	5494819	0.008633204	1078313.8
## 9	9	191802101	0.6438754	121323064	5447384	0.008353271	961052.4
## 10	10	190734044	0.6478400	120350997	5522235	0.008699768	1035524.4
## 11	11	190176874	0.6498956	119592895	5569027	0.008969923	1041122.9
## 12	12	189586789	0.6520624	119844904	5670572	0.009250883	1152643.2
## 13	13	189043550	0.6540527	119405358	5720948	0.009331370	1224627.9
## 14	14	188833635	0.6548209	119216157	5741006	0.009398136	1226570.1
## 15	15	188635164	0.6555458	119113982	5813140	0.009615091	1226816.5
## 16	16	188427962	0.6563037	118836075	5841570	0.009779370	1225391.8
## 17	17	188381716	0.6564717	118850989	5837784	0.009779374	1219600.3
## 18	18	188354341	0.6565718	118734190	5845501	0.009779065	1224907.8
## 19	19	188347708	0.6565947	118769638	5841846	0.009778721	1221526.2
## 20	20	188340273	0.6566225	118745450	5832668	0.009756786	1234081.0
## 21	21	188321979	0.6566908	118720513	5817163	0.009693742	1232452.6
## 22	22	188325117	0.6566793	118720657	5815158	0.009685612	1233727.1
## 23	23	188324643	0.6566810	118717131	5814604	0.009682559	1234709.1

Two variable we exclude is transaction_date and front_door_structure_mixed

```
#Choose the best tuned model
summary(step.model_f$finalModel)
```

```

## Subset selection object
## 23 Variables (and intercept)
##
## Forced in Forced out
## city FALSE FALSE
## transaction_year_month FALSE FALSE
## transaction_date FALSE FALSE
## year_of_completion FALSE FALSE
## exclusive_use_area FALSE FALSE
## floor FALSE FALSE
## latitude FALSE FALSE
## longitude FALSE FALSE
## address_by_law FALSE FALSE
## total_parking_capacity_in_site FALSE FALSE
## total_household_count_in_sites FALSE FALSE
## apartment_building_count_in_sites FALSE FALSE
## tallest_building_in_sites FALSE FALSE
## lowest_building_in_sites FALSE FALSE
## heat_fuel FALSE FALSE
## supply_area FALSE FALSE
## total_household_count_of_area_type FALSE FALSE
## room_count FALSE FALSE
## bathroom_count FALSE FALSE
## heat_type_central FALSE FALSE
## heat_type_district FALSE FALSE
## front_door_structure_corridor FALSE FALSE
## front_door_structure_mixed FALSE FALSE
## 1 subsets of each size up to 21
## Selection Algorithm: forward
## city transaction_year_month transaction_date year_of_completion
## 1 ( 1 ) " " " " " "
## 2 ( 1 ) "*" " " " "
## 3 ( 1 ) "*" " " " "
## 4 ( 1 ) "*" "*" " "
## 5 ( 1 ) "*" "*" " "
## 6 ( 1 ) "*" "*" " "
## 7 ( 1 ) "*" "*" " "
## 8 ( 1 ) "*" "*" " "
## 9 ( 1 ) "*" "*" " "
## 10 ( 1 ) "*" "*" " "
## 11 ( 1 ) "*" "*" " "
## 12 ( 1 ) "*" "*" "*"
## 13 ( 1 ) "*" "*" "*"
## 14 ( 1 ) "*" "*" "*"
## 15 ( 1 ) "*" "*" "*"
## 16 ( 1 ) "*" "*" "*"
## 17 ( 1 ) "*" "*" "*"
## 18 ( 1 ) "*" "*" "*"
## 19 ( 1 ) "*" "*" "*"
## 20 ( 1 ) "*" "*" "*"
## 21 ( 1 ) "*" "*" "*"
## exclusive_use_area floor latitude longitude address_by_law
## 1 ( 1 ) " " " " " "
## 2 ( 1 ) " " " " " "

```

## 3	(1)	" "	" "	" "	" "	" "
## 4	(1)	" "	" "	" "	" "	" "
## 5	(1)	" "	" "	"*"	" "	" "
## 6	(1)	" "	" "	"*"	" "	" "
## 7	(1)	" "	" "	"*"	"*"	" "
## 8	(1)	" "	" "	"*"	"*"	" "
## 9	(1)	" "	" "	"*"	"*"	" "
## 10	(1)	" "	" "	"*"	"*"	"*"
## 11	(1)	" "	" "	"*"	"*"	"*"
## 12	(1)	" "	" "	"*"	"*"	"*"
## 13	(1)	" "	"*"	"*"	"*"	"*"
## 14	(1)	" "	"*"	"*"	"*"	"*"
## 15	(1)	" "	"*"	"*"	"*"	"*"
## 16	(1)	"*"	"*"	"*"	"*"	"*"
## 17	(1)	"*"	"*"	"*"	"*"	"*"
## 18	(1)	"*"	"*"	"*"	"*"	"*"
## 19	(1)	"*"	"*"	"*"	"*"	"*"
## 20	(1)	"*"	"*"	"*"	"*"	"*"
## 21	(1)	"*"	"*"	"*"	"*"	"*"
##		total_parking_capacity_in_site	total_household_count_in_sites			
## 1	(1)	" "	" "			
## 2	(1)	" "	" "			
## 3	(1)	" "	" "			
## 4	(1)	" "	" "			
## 5	(1)	" "	" "			
## 6	(1)	" "	" "			
## 7	(1)	" "	" "			
## 8	(1)	" "	" "			
## 9	(1)	" "	"*"			
## 10	(1)	" "	"*"			
## 11	(1)	"*"	"*"			
## 12	(1)	"*"	"*"			
## 13	(1)	"*"	"*"			
## 14	(1)	"*"	"*"			
## 15	(1)	"*"	"*"			
## 16	(1)	"*"	"*"			
## 17	(1)	"*"	"*"			
## 18	(1)	"*"	"*"			
## 19	(1)	"*"	"*"			
## 20	(1)	"*"	"*"			
## 21	(1)	"*"	"*"			
##		apartment_building_count_in_sites	tallest_building_in_sites			
## 1	(1)	" "	" "			
## 2	(1)	" "	" "			
## 3	(1)	"*"	" "			
## 4	(1)	"*"	" "			
## 5	(1)	"*"	" "			
## 6	(1)	"*"	" "			
## 7	(1)	"*"	" "			
## 8	(1)	"*"	" "			
## 9	(1)	"*"	" "			
## 10	(1)	"*"	" "			
## 11	(1)	"*"	" "			
## 12	(1)	"*"	" "			

```

## 13 ( 1 ) "*" " "
## 14 ( 1 ) "*" "*"
## 15 ( 1 ) "*" "*"
## 16 ( 1 ) "*" "*"
## 17 ( 1 ) "*" "*"
## 18 ( 1 ) "*" "*"
## 19 ( 1 ) "*" "*"
## 20 ( 1 ) "*" "*"
## 21 ( 1 ) "*" "*"
## lowest_building_in_sites heat_fuel supply_area
## 1 ( 1 ) " " " " "*"
## 2 ( 1 ) " " " " "*"
## 3 ( 1 ) " " " " "*"
## 4 ( 1 ) " " " " "*"
## 5 ( 1 ) " " " " "*"
## 6 ( 1 ) " " " " "*"
## 7 ( 1 ) " " " " "*"
## 8 ( 1 ) "*" " " "*"
## 9 ( 1 ) "*" " " "*"
## 10 ( 1 ) "*" " " "*"
## 11 ( 1 ) "*" " " "*"
## 12 ( 1 ) "*" " " "*"
## 13 ( 1 ) "*" " " "*"
## 14 ( 1 ) "*" " " "*"
## 15 ( 1 ) "*" "*" "*"
## 16 ( 1 ) "*" "*" "*"
## 17 ( 1 ) "*" "*" "*"
## 18 ( 1 ) "*" "*" "*"
## 19 ( 1 ) "*" "*" "*"
## 20 ( 1 ) "*" "*" "*"
## 21 ( 1 ) "*" "*" "*"
## total_household_count_of_area_type room_count bathroom_count
## 1 ( 1 ) " " " " " "
## 2 ( 1 ) " " " " " "
## 3 ( 1 ) " " " " " "
## 4 ( 1 ) " " " " " "
## 5 ( 1 ) " " " " " "
## 6 ( 1 ) " " " " " "
## 7 ( 1 ) " " " " " "
## 8 ( 1 ) " " " " " "
## 9 ( 1 ) " " " " " "
## 10 ( 1 ) " " " " " "
## 11 ( 1 ) " " " " " "
## 12 ( 1 ) " " " " " "
## 13 ( 1 ) " " " " " "
## 14 ( 1 ) " " " " " "
## 15 ( 1 ) " " " " " "
## 16 ( 1 ) " " " " " "
## 17 ( 1 ) "*" " " " "
## 18 ( 1 ) "*" " " " "
## 19 ( 1 ) "*" "*" " "
## 20 ( 1 ) "*" "*" "*"
## 21 ( 1 ) "*" "*" "*"
## heat_type_central heat_type_district front_door_structure_corridor

```

```

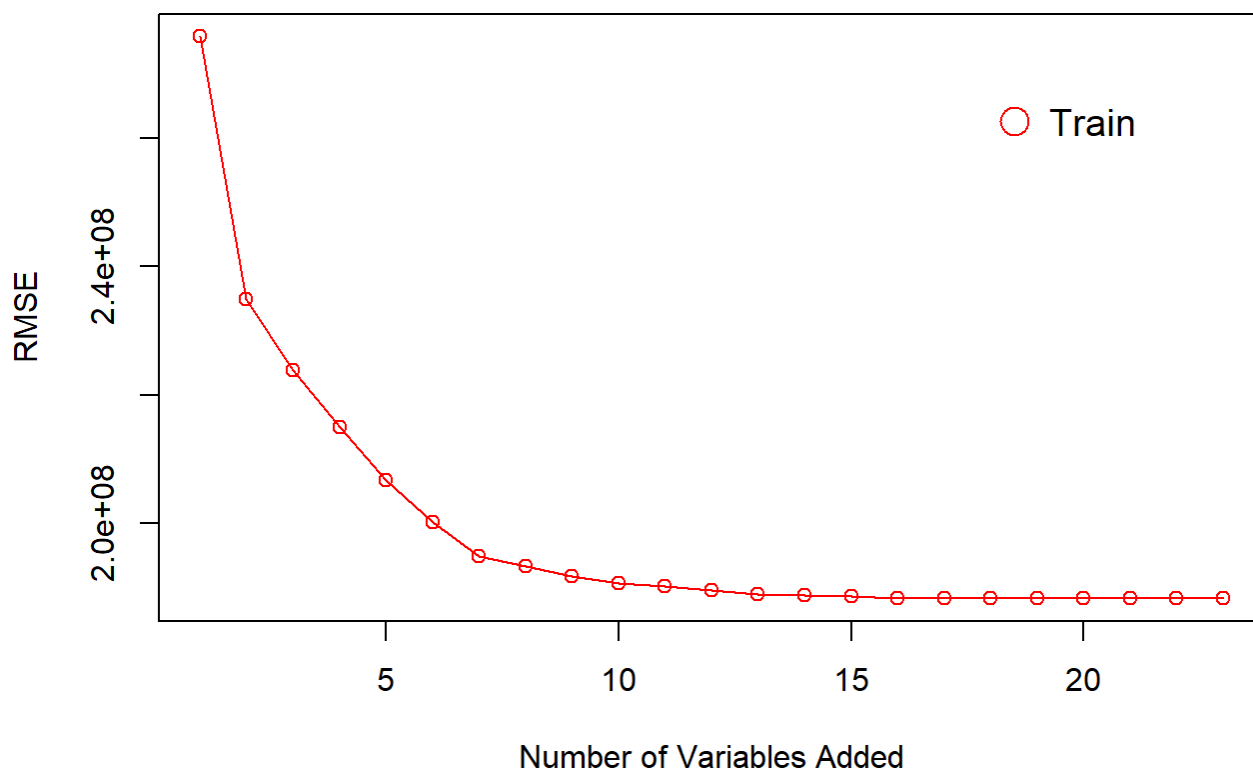
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " " "
## 5 ( 1 ) " " " "
## 6 ( 1 ) " " "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) " " "*"
## 9 ( 1 ) " " "*"
## 10 ( 1 ) " " "*"
## 11 ( 1 ) " " "*"
## 12 ( 1 ) " " "*"
## 13 ( 1 ) " " "*"
## 14 ( 1 ) " " "*"
## 15 ( 1 ) " " "*"
## 16 ( 1 ) " " "*"
## 17 ( 1 ) " " "*"
## 18 ( 1 ) " " "*"
## 19 ( 1 ) " " "*"
## 20 ( 1 ) " " "*"
## 21 ( 1 ) "*" "*" "*"

## front_door_structure_mixed
## 1 ( 1 ) " "
## 2 ( 1 ) " "
## 3 ( 1 ) " "
## 4 ( 1 ) " "
## 5 ( 1 ) " "
## 6 ( 1 ) " "
## 7 ( 1 ) " "
## 8 ( 1 ) " "
## 9 ( 1 ) " "
## 10 ( 1 ) " "
## 11 ( 1 ) " "
## 12 ( 1 ) " "
## 13 ( 1 ) " "
## 14 ( 1 ) " "
## 15 ( 1 ) " "
## 16 ( 1 ) " "
## 17 ( 1 ) " "
## 18 ( 1 ) " "
## 19 ( 1 ) " "
## 20 ( 1 ) " "
## 21 ( 1 ) " "

```

```
x1f <- step.model_f$results[,2]
plot(x1f,type = "o",col = "red",xlab = "Number of Variables Added", ylab = "RMSE",
     main = "Forward Train RMSE")
legend("topright",
      legend = c("Train"),
      col = c("red"),
      pch = c(1,1),
      bty = "n",
      pt.cex = 2,
      cex = 1.2,
      text.col = "black",
      horiz = F ,
      inset = c(0.1, 0.1))
```

Forward Train RMSE



Backward selection

In this model, we came to the same conclusion that the most optimal model has 21 variable with the lowest RMSE of 188,321,979 Korean Won.

```

set.seed(810)
#Backward train
# Set up repeated k-fold cross-validation
train.control <- trainControl(method = "cv", number = 10)
# Train the model
step.model_b <- train(transaction_real_price ~., data = dd_train,
                      method = "leapBackward",
                      tuneGrid = data.frame(nvmax = 1:23),
                      trControl = train.control
                      )
step.model_b$results

```

##	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	275820794	0.2633316	183550247	5271468	0.012969232	1810020.1
## 2	2	234909141	0.4656794	153777481	5215660	0.009970077	1767563.1
## 3	3	225983655	0.5055657	147442830	5086690	0.008856248	1418133.3
## 4	4	216077790	0.5479695	141673953	5201928	0.009422754	1347680.6
## 5	5	206243197	0.5882408	132030639	5346512	0.008208107	1244955.9
## 6	6	200214993	0.6119412	126968771	5216558	0.008366577	1771089.2
## 7	7	194816961	0.6325771	123768489	5318325	0.008077787	1169235.0
## 8	8	193088790	0.6390701	122727058	5497882	0.008709827	1068180.3
## 9	9	191591793	0.6446538	121386085	5458523	0.008446345	973728.1
## 10	10	190617456	0.6482669	120455377	5536638	0.008785837	1044868.6
## 11	11	190082271	0.6502400	119726562	5583167	0.009057233	1054663.3
## 12	12	189478964	0.6524544	119972543	5689114	0.009350330	1177436.5
## 13	13	188935542	0.6544445	119543844	5737249	0.009421530	1241138.6
## 14	14	188699814	0.6553071	119325857	5759809	0.009505695	1226964.7
## 15	15	188635164	0.6555458	119113982	5813140	0.009615091	1226816.5
## 16	16	188427962	0.6563037	118836075	5841570	0.009779370	1225391.8
## 17	17	188381716	0.6564717	118850989	5837784	0.009779374	1219600.3
## 18	18	188354341	0.6565718	118734190	5845501	0.009779065	1224907.8
## 19	19	188347708	0.6565947	118769638	5841846	0.009778721	1221526.2
## 20	20	188340273	0.6566225	118745450	5832668	0.009756786	1234081.0
## 21	21	188321979	0.6566908	118720513	5817163	0.009693742	1232452.6
## 22	22	188325117	0.6566793	118720657	5815158	0.009685612	1233727.1
## 23	23	188324643	0.6566810	118717131	5814604	0.009682559	1234709.1

Two variable we exclude is transaction_date and front_door_structure_mixed

```

#Choose the best tuned model
summary(step.model_b$finalModel)

```

```

## Subset selection object
## 23 Variables (and intercept)
##
## Forced in Forced out
## city FALSE FALSE
## transaction_year_month FALSE FALSE
## transaction_date FALSE FALSE
## year_of_completion FALSE FALSE
## exclusive_use_area FALSE FALSE
## floor FALSE FALSE
## latitude FALSE FALSE
## longitude FALSE FALSE
## address_by_law FALSE FALSE
## total_parking_capacity_in_site FALSE FALSE
## total_household_count_in_sites FALSE FALSE
## apartment_building_count_in_sites FALSE FALSE
## tallest_building_in_sites FALSE FALSE
## lowest_building_in_sites FALSE FALSE
## heat_fuel FALSE FALSE
## supply_area FALSE FALSE
## total_household_count_of_area_type FALSE FALSE
## room_count FALSE FALSE
## bathroom_count FALSE FALSE
## heat_type_central FALSE FALSE
## heat_type_district FALSE FALSE
## front_door_structure_corridor FALSE FALSE
## front_door_structure_mixed FALSE FALSE
## 1 subsets of each size up to 21
## Selection Algorithm: backward
## city transaction_year_month transaction_date year_of_completion
## 1 ( 1 ) " " " " " " " "
## 2 ( 1 ) "*" " " " " " "
## 3 ( 1 ) "*" " " " " " "
## 4 ( 1 ) "*" " " " " " "
## 5 ( 1 ) "*" "*" " " " "
## 6 ( 1 ) "*" "*" " " " "
## 7 ( 1 ) "*" "*" " " " "
## 8 ( 1 ) "*" "*" " " " "
## 9 ( 1 ) "*" "*" " " " "
## 10 ( 1 ) "*" "*" " " " "
## 11 ( 1 ) "*" "*" " " " "
## 12 ( 1 ) "*" "*" " " "*"
## 13 ( 1 ) "*" "*" " " "*"
## 14 ( 1 ) "*" "*" " " "*"
## 15 ( 1 ) "*" "*" " " "*"
## 16 ( 1 ) "*" "*" " " "*"
## 17 ( 1 ) "*" "*" " " "*"
## 18 ( 1 ) "*" "*" " " "*"
## 19 ( 1 ) "*" "*" " " "*"
## 20 ( 1 ) "*" "*" " " "*"
## 21 ( 1 ) "*" "*" " " "*"
##
## exclusive_use_area floor latitude longitude address_by_law
## 1 ( 1 ) " " " " " " " "
## 2 ( 1 ) " " " " " " " "

```



```

## 3 ( 1 ) " " " " "*" " " " "
## 4 ( 1 ) " " " " "*" " " " "
## 5 ( 1 ) " " " " "*" " " " "
## 6 ( 1 ) " " " " "*" " " " "
## 7 ( 1 ) " " " " "*" "*" " "
## 8 ( 1 ) " " " " "*" "*" " "
## 9 ( 1 ) " " " " "*" "*" " "
## 10 ( 1 ) " " " " "*" "*" "*"
## 11 ( 1 ) " " " " "*" "*" "*"
## 12 ( 1 ) " " " " "*" "*" "*"
## 13 ( 1 ) " " "*" "*" "*" "*"
## 14 ( 1 ) " " "*" "*" "*" "*"
## 15 ( 1 ) " " "*" "*" "*" "*"
## 16 ( 1 ) "*" "*" "*" "*" "*"
## 17 ( 1 ) "*" "*" "*" "*" "*"
## 18 ( 1 ) "*" "*" "*" "*" "*"
## 19 ( 1 ) "*" "*" "*" "*" "*"
## 20 ( 1 ) "*" "*" "*" "*" "*"
## 21 ( 1 ) "*" "*" "*" "*" "*"

## total_parking_capacity_in_site total_household_count_in_sites
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " " "
## 5 ( 1 ) " " " "
## 6 ( 1 ) " " " "
## 7 ( 1 ) " " " "
## 8 ( 1 ) " " " "
## 9 ( 1 ) " " "*"
## 10 ( 1 ) " " "*"
## 11 ( 1 ) "*" "*"
## 12 ( 1 ) "*" "*"
## 13 ( 1 ) "*" "*"
## 14 ( 1 ) "*" "*"
## 15 ( 1 ) "*" "*"
## 16 ( 1 ) "*" "*"
## 17 ( 1 ) "*" "*"
## 18 ( 1 ) "*" "*"
## 19 ( 1 ) "*" "*"
## 20 ( 1 ) "*" "*"
## 21 ( 1 ) "*" "*"

## apartment_building_count_in_sites tallest_building_in_sites
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " " "
## 5 ( 1 ) " " " "
## 6 ( 1 ) "*" " "
## 7 ( 1 ) "*" " "
## 8 ( 1 ) "*" " "
## 9 ( 1 ) "*" " "
## 10 ( 1 ) "*" " "
## 11 ( 1 ) "*" " "
## 12 ( 1 ) "*" " "

```

```

## 13 ( 1 ) "*"
## 14 ( 1 ) "*"
## 15 ( 1 ) "*"
## 16 ( 1 ) "*"
## 17 ( 1 ) "*"
## 18 ( 1 ) "*"
## 19 ( 1 ) "*"
## 20 ( 1 ) "*"
## 21 ( 1 ) "*"
##
## lowest_building_in_sites heat_fuel supply_area
## 1 ( 1 ) " " " " "*"
## 2 ( 1 ) " " " " "*"
## 3 ( 1 ) " " " " "*"
## 4 ( 1 ) " " " " "*"
## 5 ( 1 ) " " " " "*"
## 6 ( 1 ) " " " " "*"
## 7 ( 1 ) " " " " "*"
## 8 ( 1 ) "*" " " "*"
## 9 ( 1 ) "*" " " "*"
## 10 ( 1 ) "*" " " "*"
## 11 ( 1 ) "*" " " "*"
## 12 ( 1 ) "*" " " "*"
## 13 ( 1 ) "*" " " "*"
## 14 ( 1 ) "*" " " "*"
## 15 ( 1 ) "*" "*" "*"
## 16 ( 1 ) "*" "*" "*"
## 17 ( 1 ) "*" "*" "*"
## 18 ( 1 ) "*" "*" "*"
## 19 ( 1 ) "*" "*" "*"
## 20 ( 1 ) "*" "*" "*"
## 21 ( 1 ) "*" "*" "*"
##
## total_household_count_of_area_type room_count bathroom_count
## 1 ( 1 ) " " " " " "
## 2 ( 1 ) " " " " " "
## 3 ( 1 ) " " " " " "
## 4 ( 1 ) " " " " " "
## 5 ( 1 ) " " " " " "
## 6 ( 1 ) " " " " " "
## 7 ( 1 ) " " " " " "
## 8 ( 1 ) " " " " " "
## 9 ( 1 ) " " " " " "
## 10 ( 1 ) " " " " " "
## 11 ( 1 ) " " " " " "
## 12 ( 1 ) " " " " " "
## 13 ( 1 ) " " " " " "
## 14 ( 1 ) " " " " " "
## 15 ( 1 ) " " " " " "
## 16 ( 1 ) " " " " " "
## 17 ( 1 ) "*" " " " "
## 18 ( 1 ) "*" " " " "
## 19 ( 1 ) "*" "*" " "
## 20 ( 1 ) "*" "*" "*"
## 21 ( 1 ) "*" "*" "*"
##
## heat_type_central heat_type_district front_door_structure_corridor

```

```

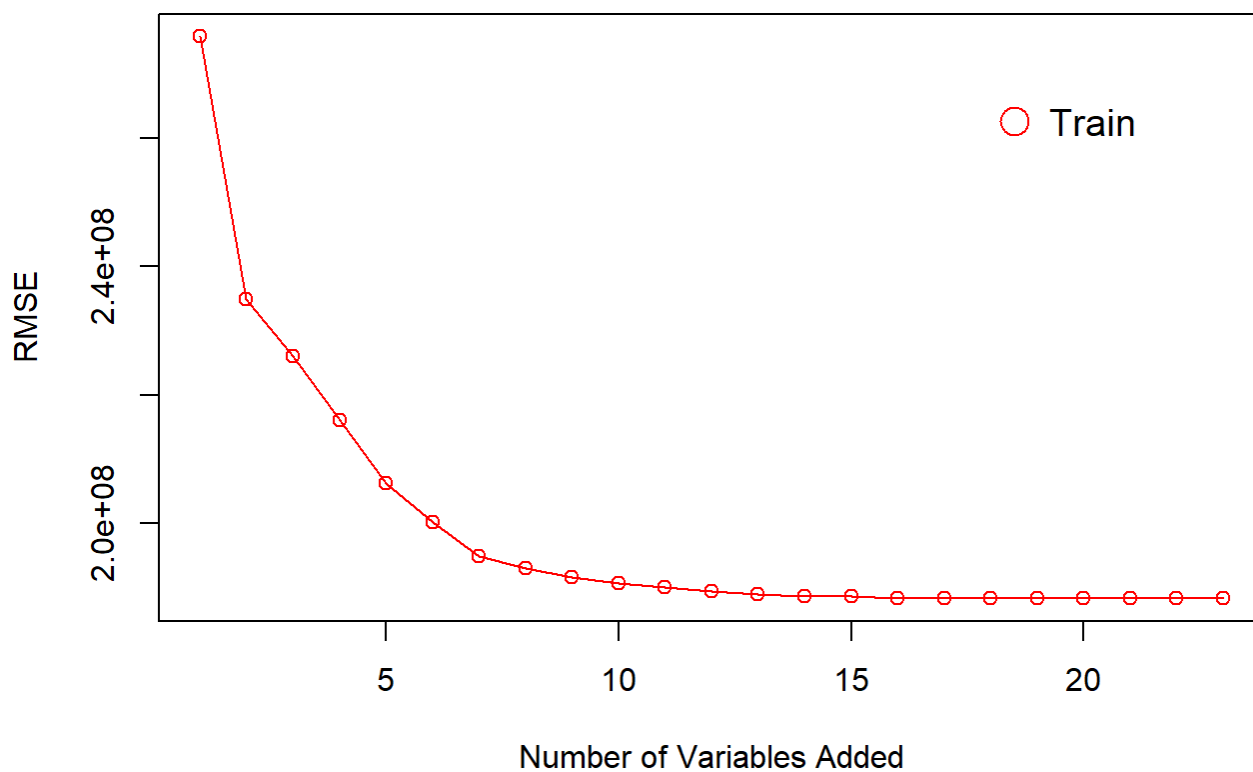
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " "*"
## 5 ( 1 ) " " "*"
## 6 ( 1 ) " " "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) " " "*"
## 9 ( 1 ) " " "*"
## 10 ( 1 ) " " "*"
## 11 ( 1 ) " " "*"
## 12 ( 1 ) " " "*"
## 13 ( 1 ) " " "*"
## 14 ( 1 ) " " "*"
## 15 ( 1 ) " " "*"
## 16 ( 1 ) " " "*"
## 17 ( 1 ) " " "*"
## 18 ( 1 ) " " "*"
## 19 ( 1 ) " " "*"
## 20 ( 1 ) " " "*"
## 21 ( 1 ) "*" "*" "*"

## front_door_structure_mixed
## 1 ( 1 ) " "
## 2 ( 1 ) " "
## 3 ( 1 ) " "
## 4 ( 1 ) " "
## 5 ( 1 ) " "
## 6 ( 1 ) " "
## 7 ( 1 ) " "
## 8 ( 1 ) " "
## 9 ( 1 ) " "
## 10 ( 1 ) " "
## 11 ( 1 ) " "
## 12 ( 1 ) " "
## 13 ( 1 ) " "
## 14 ( 1 ) " "
## 15 ( 1 ) " "
## 16 ( 1 ) " "
## 17 ( 1 ) " "
## 18 ( 1 ) " "
## 19 ( 1 ) " "
## 20 ( 1 ) " "
## 21 ( 1 ) " "

```

```
x1 <- step.model_b$results[,2]
plot(x1,type = "o",col = "red",xlab = "Number of Variables Added", ylab = "RMSE",
     main = "Backward Train RMSE")
legend("topright",
      legend = c("Train"),
      col = c("red"),
      pch = c(1,1),
      bty = "n",
      pt.cex = 2,
      cex = 1.2,
      text.col = "black",
      horiz = F ,
      inset = c(0.1, 0.1))
```

Backward Train RMSE



Since the model chosen by Forward/Backward is the same containing the same variables, let's fit it to calculate the final test RMSE. So, Forward/Backward Selection Regression Test RMSE is 188,550,321 Korean Won.

```
# Fit the model chosen by forward and backward
set.seed(810)
train.control <- trainControl(method = "cv", number = 10)
# Train the model
fit_fb <- train(transaction_real_price ~ city + transaction_year_month + year_of_completion + exclusive_use_area + floor + longitude + latitude + address_by_law + total_parking_capacity_in_site + total_household_count_in_sites + apartment_building_count_in_sites + tallest_building_in_sites + lowest_building_in_sites + heat_type_central + heat_type_district + heat_fuel + supply_area + total_household_count_of_area_type + room_count + bathroom_count + front_door_structure_corridor, data=dd_train, method = "lm",
                      trControl = train.control)
# Summarize the results
print(fit_fb)
```

```
## Linear Regression
##
## 128000 samples
## 21 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 115201, 115201, 115199, 115201, 115199, 115199, ...
## Resampling results:
##
## RMSE      Rsquared   MAE
## 188321979  0.6566908  118720513
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

```
yhat_test_fb <- predict(fit_fb, dd_test)
mse_test_fb <- mean((y_test - yhat_test_fb)**2)
paste("Forward/Backward Selection Regression Test RMSE", sqrt(mse_test_fb))
```

```
## [1] "Forward/Backward Selection Regression Test RMSE 188550321.251422"
```

Setup for Ridge and Lasso Regression

We created a new formula that includes all 23 predictors, and split out test and train datasets

```
# added all of the variables to the formula so that we can have 24 predictors
f2 <- as.formula(transaction_real_price ~ city + transaction_year_month + transaction_date + year_of_completion + exclusive_use_area + floor + longitude + latitude + address_by_law + total_parking_capacity_in_site + total_household_count_in_sites + apartment_building_count_in_sites + tallest_building_in_sites + lowest_building_in_sites + heat_type_central + heat_type_district + heat_fuel + supply_area + total_household_count_of_area_type + room_count + bathroom_count + front_door_structure_mixed + front_door_structure_corridor)

x1_train_sample <- model.matrix(f2, dd_train)[, -1]
x1_test <- model.matrix(f2, dd_test)[, -1]
```

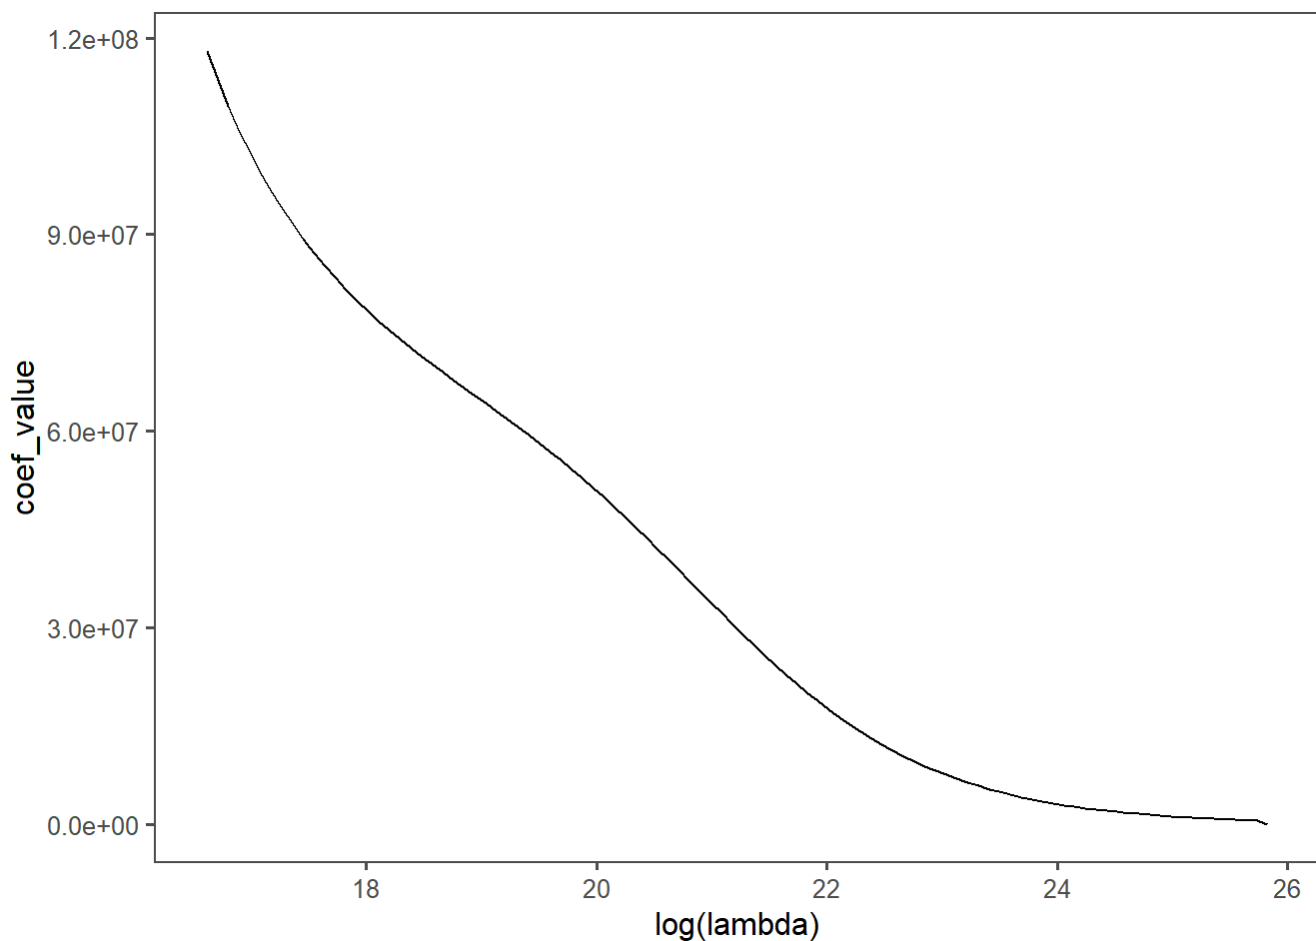
Ridge Regression

We used cross validation to identify the best lambda value. We can see that this model actually performs slightly worse compared to linear regression. The graph reveals that as the training ran, our coefficients converged towards zero.

```
# fit the ridge regression using the cross validation data
fit.ridge <- cv.glmnet(x1_train_sample, y_train, alpha = 0)

# make predictions using fitted model
ridge.coef <- predict(fit.ridge,
  type = "coefficients",
  s = fit.ridge$lambda)
to_plot <- data.table(
  lambda = fit.ridge$lambda,
  coef_value = ridge.coef[2, ]
)

# plot the coefficient values for different values of lambda
ggplot(to_plot, aes(log(lambda), coef_value)) +
  geom_line() +
  theme_few()
```



```
# MSE for train
yhat.train.ridge <- predict(fit.ridge, x1_train_sample, s = fit.ridge$lambda.min)
mse.train.ridge <- mean((y_train - yhat.train.ridge)^2)

# MSE to test
yhat.test.ridge <- predict(fit.ridge, x1_test, s = fit.ridge$lambda.min)
mse.test.ridge <- mean((y_test - yhat.test.ridge)^2)

cat("Train RMSE: ",sqrt(mse.train.ridge))
```

```
## Train RMSE: 201423543
```

```
cat(" Test RMSE: ",sqrt(mse.test.ridge))
```

```
## Test RMSE: 201931073
```

```
cat(" Best Lambda: ", fit.ridge$lambda.min)
```

```
## Best Lambda: 16491016
```

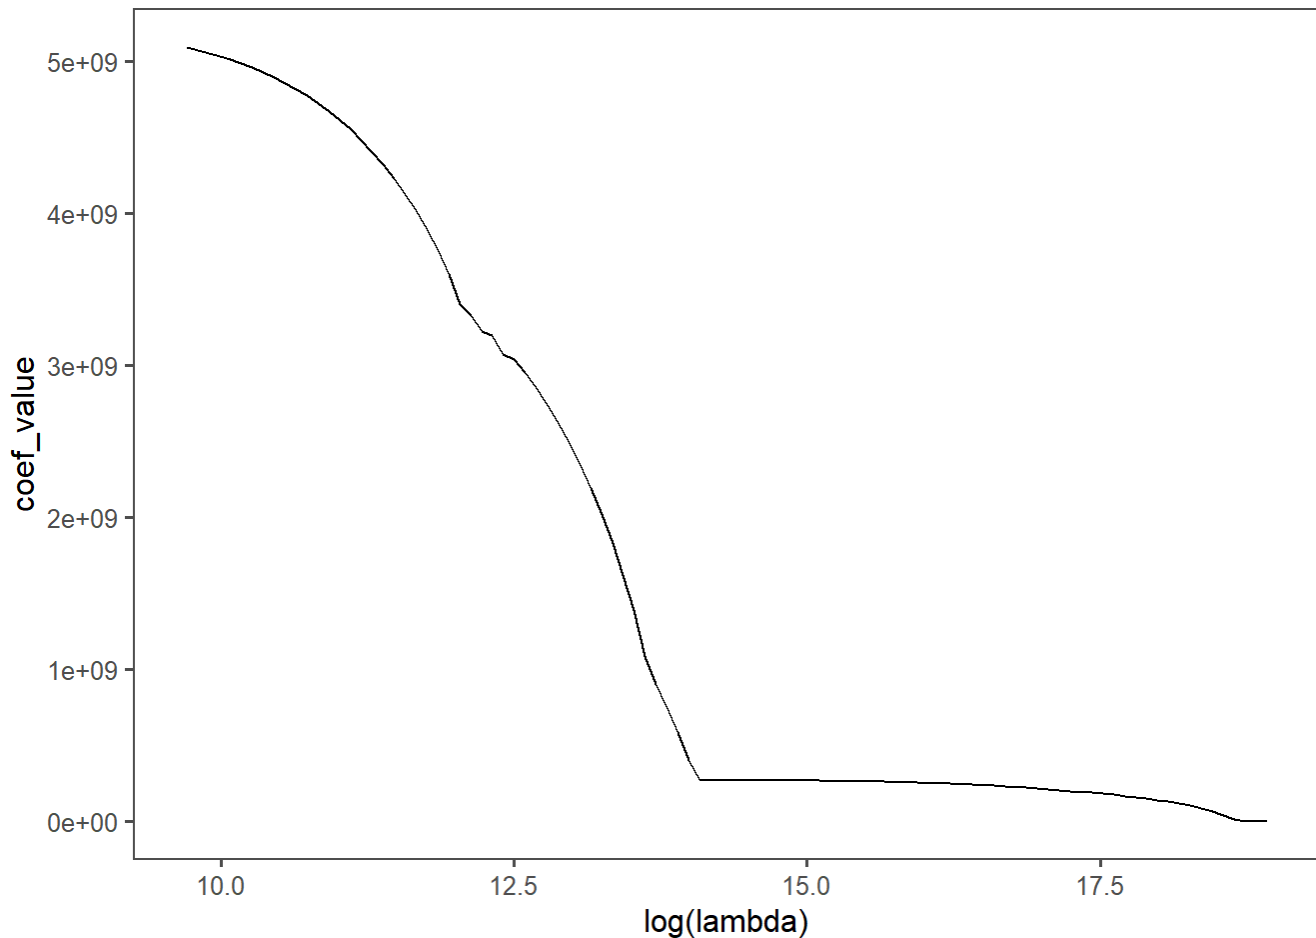
Lasso Regression

In this model, we again used cross validation to find the optimal lambda value. In the output below, we can see the coefficient value associated with each of our predictors. None of the coefficient values are zero, increasing confidence that we are not overfitting, and that all of our features are contributing to the model.

```
fit.lasso <- cv.glmnet(x1_train_sample, y_train, alpha = 1)

# predict based on most optimal lambda found above
lasso.coef <- predict(fit.lasso,
  type = "coefficients",
  s = fit.lasso$lambda)
to_plot <- data.table(
  lambda = fit.lasso$lambda,
  coef_value = lasso.coef[2, ]
)

# plot the coefficient values for different values of lambda
ggplot(to_plot, aes(log(lambda), coef_value)) +
  geom_line() +
  theme_few()
```



```
yhat.train.lasso <- predict(fit.lasso, x1_train_sample, s = fit.lasso$lambda.min)
mse.train.lasso <- mean((y_train - yhat.train.lasso)^2)
```

```
yhat.test.lasso <- predict(fit.lasso, x1_test, s = fit.lasso$lambda.min)
mse.test.lasso <- mean((y_test - yhat.test.lasso)^2)
```

```
cat("Train RMSE: ",sqrt(mse.train.lasso))
```

```
## Train RMSE: 188454633
```

```
cat(" Test RMSE: ",sqrt(mse.test.lasso))
```

```
## Test RMSE: 188678874
```

```
cat(" Best Lambda: ", fit.lasso$lambda.min)
```

```
## Best Lambda: 16491.02
```

Decision Tree

In our decision tree model, we use `row_index` to get the same observations used in the previous models.

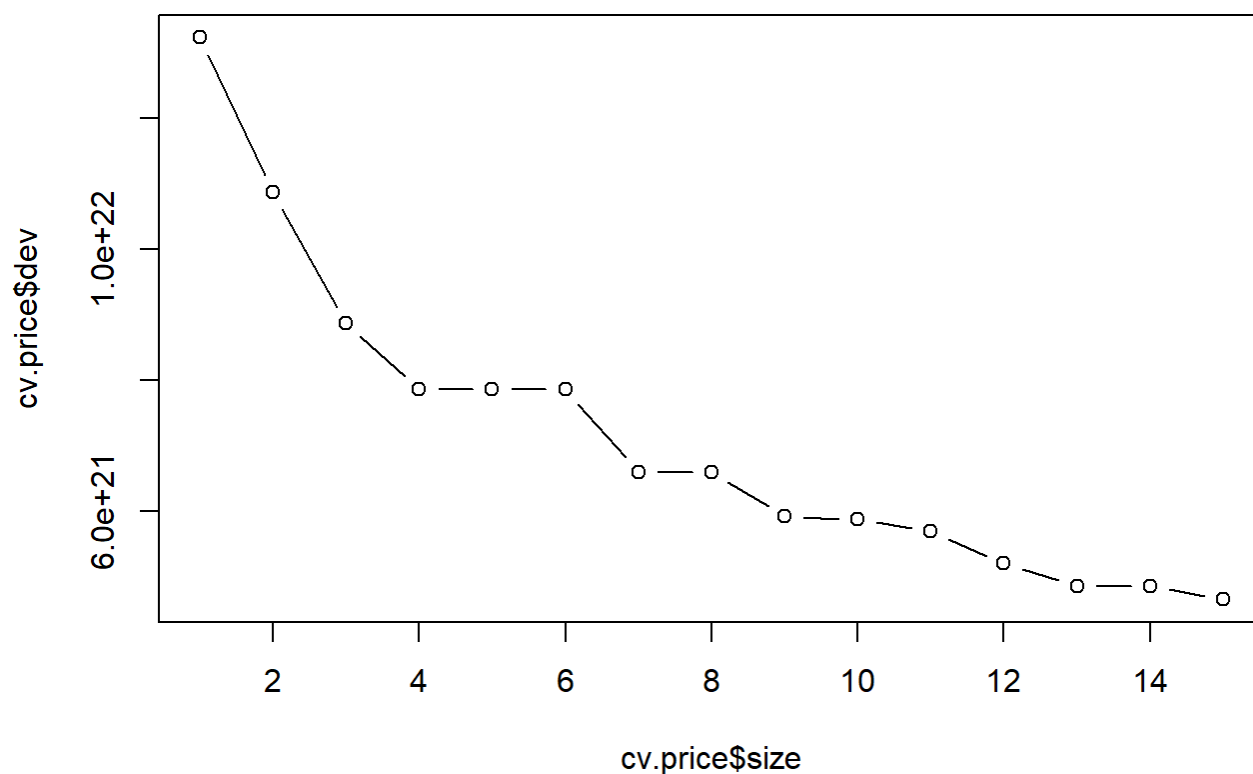

```
tree.price = tree(transaction_real_price ~ . , Price, subset = row_index)

summary(tree.price)
```

```
##
## Regression tree:
## tree(formula = transaction_real_price ~ ., data = Price, subset = row_index)
## Variables actually used in tree construction:
## [1] "supply_area"          "latitude"
## [3] "address_by_law"       "exclusive_use_area"
## [5] "heat_fuel"            "longitude"
## [7] "tallest_building_in_sites" "apartment_building_count_in_sites"
## [9] "transaction_year_month"
## Number of terminal nodes: 15
## Residual mean deviance: 3.519e+16 = 4.504e+21 / 128000
## Distribution of residuals:
##      Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
## -1.691e+09 -9.586e+07 -2.407e+07  0.000e+00  6.814e+07  3.883e+09
```

We can see that we don't have to prune that tree because the largest tree (size = 14), has the lowest cross validation error.

```
# Prune Tree
cv.price = cv.tree(tree.price)
plot(cv.price$size, cv.price$dev, type = "b")
```



cv.price

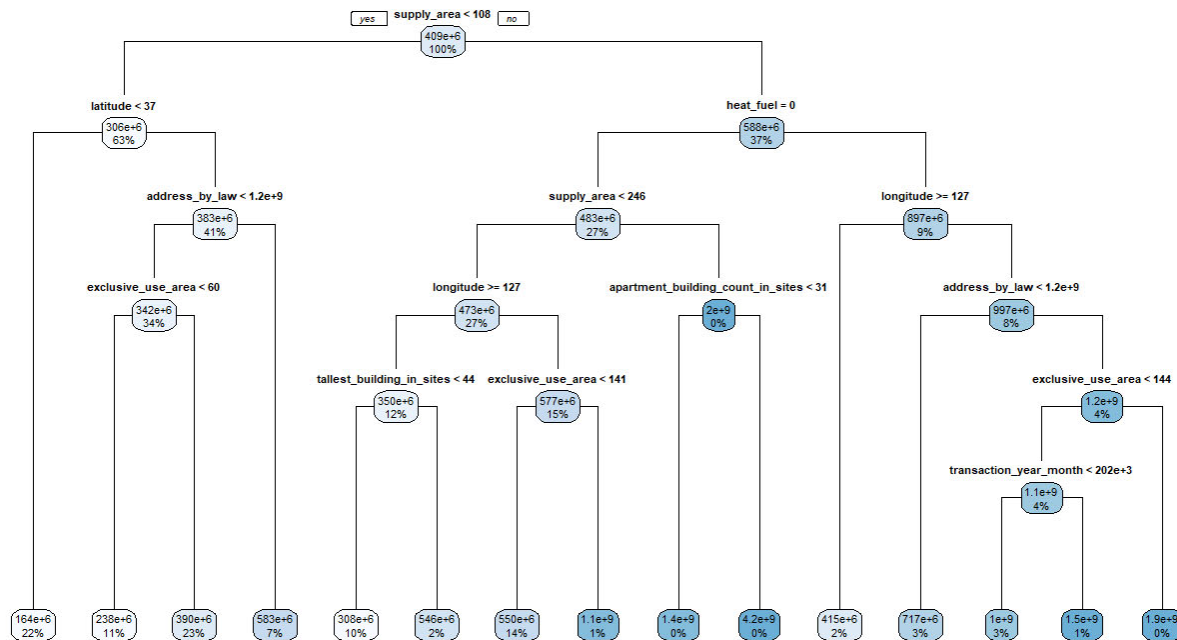
```
## $size
## [1] 15 14 13 12 11 10 9 8 7 6 5 4 3 2 1
##
## $dev
## [1] 4.657197e+21 4.864690e+21 4.864690e+21 5.202028e+21 5.700069e+21
## [6] 5.873078e+21 5.925532e+21 6.598284e+21 6.598284e+21 7.858926e+21
## [11] 7.858926e+21 7.858926e+21 8.862230e+21 1.086308e+22 1.322130e+22
##
## $k
## [1] -Inf 1.333288e+20 1.368796e+20 2.168168e+20 2.952076e+20
## [6] 3.052146e+20 3.356930e+20 4.303563e+20 4.439117e+20 5.320393e+20
## [11] 5.576437e+20 5.724015e+20 8.868482e+20 1.512390e+21 2.358311e+21
##
## $method
## [1] "deviance"
##
## attr(,"class")
## [1] "prune" "tree.sequence"
```

Our RMSE below is 189,891,653.318 Korean Won.

```
#Test MSE
tree.yhat = predict(tree.price,newdata=dd_test)
mean((tree.yhat - y_test)^2)
```

```
## [1] 3.692462e+16
```

```
#Lecture Method to deal with decision tree
set.seed(217)
tree.price.lec <- rpart(transaction_real_price ~ ., Price, subset = row_index,
                        control = rpart.control(cp=0.01))
rpart.plot(tree.price.lec, type = 1)
```



Bagging

```
#### Trends of Test MSE as number of data in training growing (1% to 5%)
#set.seed(217)
#test.mse = c()

#for (i in seq(1,5)) {
#  train = sample(1:nrow(Price),(nrow(Price)/100)*i)
#  tree.testy = Price[-train,transaction_real_price]
#  tree.test = Price[-train]
#  bag.price = randomForest(transaction_real_price ~ ., data = Price, subset = train, mtry = 24,
importance = TRUE)
#  yhat.bag = predict(bag.price, newdata = tree.test)
#  test.mse = c(test.mse,mean((tree.testy-yhat.bag)^2))
#}

#test.mse
```

Bagging using 5000 rows as training

We decided to use 5000 rows for compute resource reasons. We then used the remaining data to calculate the test MSE. We found that there was a significant improvement in the RMSE, which we calculate to be 101,033,509.293 Korean Won. Also, we found that when we increased the number of observations in the train dataset, our MSE went down significantly.

```
bag.price = randomForest(transaction_real_price ~ ., data = dd_train, mtry=21, importance = TRUE
)
yhat.bag.train = predict(bag.price, newdata = dd_train)
cat("RMSE train: ", sqrt(mean((y_train-yhat.bag.train)^2)))
```

```
## RMSE train: 21527264
```

```
yhat.bag = predict(bag.price, newdata = dd_test)
cat(" RMSE test: ", sqrt(mean((y_test-yhat.bag)^2)))
```

```
## RMSE test: 48105079
```

Random Forest

This process is similar to bagging; we took a sample of 5000 observations from the dataset. We then ran the random forest model on out sample, and computed an MSE. The MSE was again an improvement over some of the less flexible models earlier in the report. We calculated a test MSE of 102,026,222.12 Korean Won. Due to the fact that we are only able to take a sample of 5000 observations, the model has higher variability since the MSE can change significantly every time we run the model.

```
rf.price = randomForest(transaction_real_price ~ ., data = Price, subset = row_index, mtry = 5,
importance = TRUE)

yhat.rf.train = predict(rf.price, newdata = dd_train)
cat("RMSE train: ", sqrt(mean((y_train-yhat.rf.train)^2)))
```

```
## RMSE train: 34080272
```

```
yhat.rf = predict(rf.price, newdata = dd_test)
cat(" RMSE test: ", sqrt(mean((y_test-yhat.rf)^2)))
```

```
## RMSE test: 54993204
```

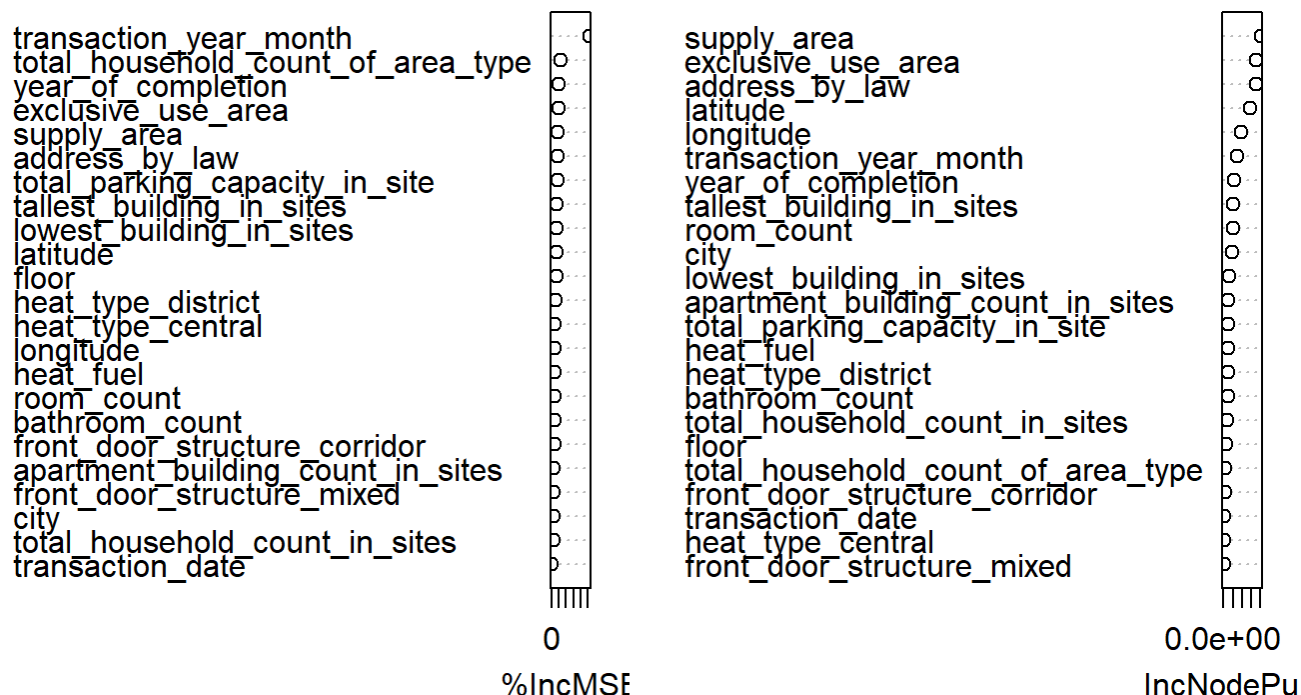
```
importance(rf.price)
```

##	%IncMSE	IncNodePurity
## city	13.296800	4.672313e+20
## transaction_year_month	261.082473	7.331360e+20
## transaction_date	2.130253	2.112215e+19
## year_of_completion	51.249919	5.665932e+20
## exclusive_use_area	47.950984	1.794052e+21
## floor	33.677676	1.163679e+20
## latitude	34.459616	1.452699e+21
## longitude	21.894668	9.760793e+20
## address_by_law	43.654317	1.771232e+21
## total_parking_capacity_in_site	40.089590	2.738914e+20
## total_household_count_in_sites	12.150153	2.007782e+20
## apartment_building_count_in_sites	16.257821	2.741756e+20
## tallest_building_in_sites	35.551012	5.369457e+20
## lowest_building_in_sites	34.851280	3.135553e+20
## heat_fuel	21.851396	2.737434e+20
## supply_area	43.904885	2.026513e+21
## total_household_count_of_area_type	62.743647	1.100054e+20
## room_count	21.126094	4.995193e+20
## bathroom_count	19.225478	2.360294e+20
## heat_type_central	22.830232	1.338015e+19
## heat_type_district	27.507456	2.674862e+20
## front_door_structure_corridor	18.231974	7.695837e+19
## front_door_structure_mixed	13.825165	8.157309e+18

The visualization below reveals that supply area is relatively more important compared to the other predictors. Exclusive use area is also an important predictor.

```
varImpPlot(rf.price)
```

rf.price



Boosting

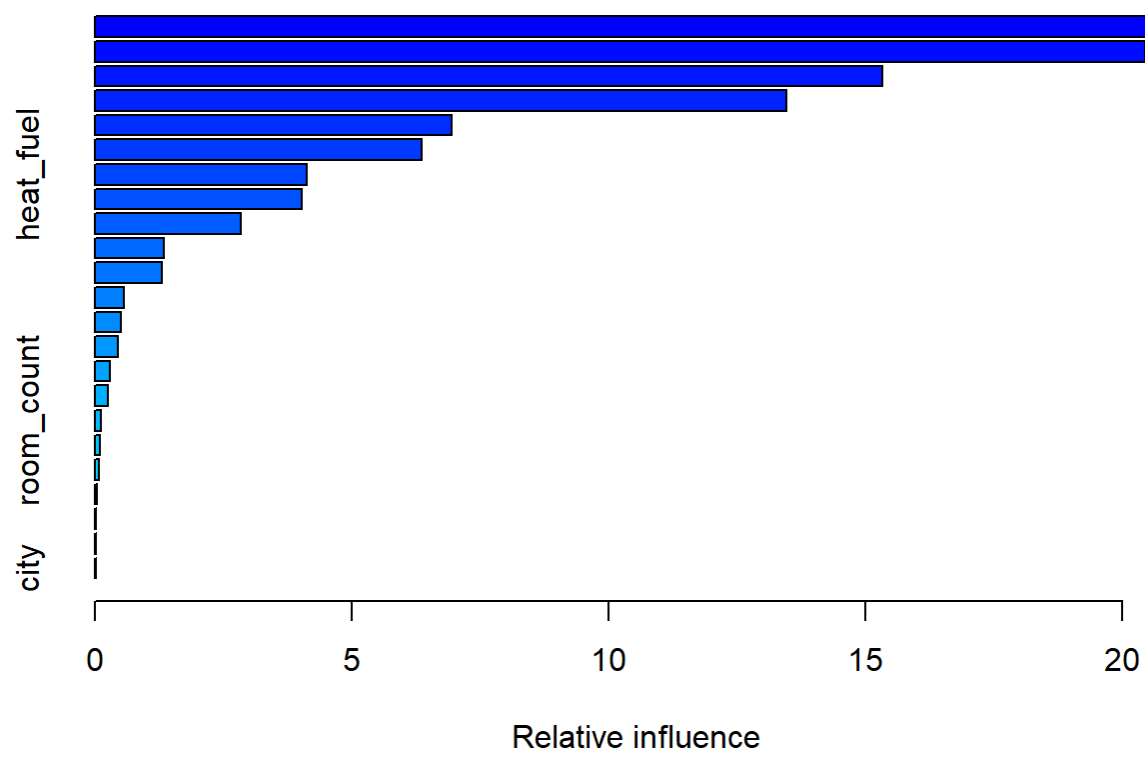
the boosting process reveals similar finding to what we saw with random forest above. Again, we can see that supply area, and exclusive use area are both relatively more important compared to the other predictors.

select a random sample of 10000 observations

```
# get a random sample of row numbers from the train dataset
train_sample_index = sample(1:nrow(dd_train),nrow(dd_train)*4/5)
# use those rows to get the actual data from the train data
train_sample = dd_train[train_sample_index,]
# get the price column from the train dataset
tree.testy_sample <- dd_test[train_sample_index, transaction_real_price]
```

In the below model, we started with just 5000 trees. We then go on to perform feature engineering, and use cross validation to find the optimal number of trees.

```
boost.price = gbm(transaction_real_price ~ ., data = train_sample, distribution = "gaussian", n.
trees = 5000, interaction.depth = 4 )
summary(boost.price)
```



```

##                                     var
## address_by_law                     address_by_law
## supply_area                        supply_area
## exclusive_use_area                 exclusive_use_area
## latitude                           latitude
## transaction_year_month             transaction_year_month
## longitude                           longitude
## heat_fuel                           heat_fuel
## year_of_completion                 year_of_completion
## tallest_building_in_sites          tallest_building_in_sites
## apartment_building_count_in_sites apartment_building_count_in_sites
## total_parking_capacity_in_site     total_parking_capacity_in_site
## lowest_building_in_sites           lowest_building_in_sites
## total_household_count_in_sites     total_household_count_in_sites
## floor                              floor
## heat_type_district                 heat_type_district
## total_household_count_of_area_type total_household_count_of_area_type
## room_count                         room_count
## front_door_structure_corridor      front_door_structure_corridor
## bathroom_count                     bathroom_count
## heat_type_central                   heat_type_central
## transaction_date                    transaction_date
## front_door_structure_mixed         front_door_structure_mixed
## city                               city
##                                   rel.inf
## address_by_law                     21.30022955
## supply_area                        20.44244313
## exclusive_use_area                 15.33593540
## latitude                           13.47260822
## transaction_year_month             6.94010123
## longitude                           6.36395481
## heat_fuel                           4.12824315
## year_of_completion                 4.03493562
## tallest_building_in_sites          2.83937368
## apartment_building_count_in_sites 1.33846569
## total_parking_capacity_in_site     1.31318606
## lowest_building_in_sites           0.56585616
## total_household_count_in_sites     0.50259046
## floor                              0.44629422
## heat_type_district                 0.30050003
## total_household_count_of_area_type 0.25592111
## room_count                         0.12440337
## front_door_structure_corridor      0.09784408
## bathroom_count                     0.08454600
## heat_type_central                   0.04972401
## transaction_date                    0.03084237
## front_door_structure_mixed         0.01854806
## city                               0.01345361

```

The results of the boosting model reveal the following RMSEs in Korean Won. This model produced the lowest error of any of the models in the report.


```
yhat.boost.train = predict(boost.price, newdata = dd_train, n.trees = 5000)
cat("RMSE train: ", sqrt(mean((dd_train$transaction_real_price - yhat.boost.train)^2)))
```

```
## RMSE train: 44452017
```

```
#Evaluate boosted tree model
```

```
yhat.boost.test = predict(boost.price, newdata = dd_test, n.trees = 5000)
cat(" RMSE test: ", sqrt(mean((dd_test$transaction_real_price - yhat.boost.test)^2)))
```

```
## RMSE test: 52281154
```

Feature Engineering

```
dd_train_engineering <- copy(dd_train)
dd_test_engineering <- copy(dd_test)
```

```
# feature engineering for train
```

```
# perform feature engineering to put the latitude and longitude in 3d space
```

```
dd_train_engineering[, x := cos(latitude) * cos(longitude)]
```

```
dd_train_engineering[, y := cos(latitude) * sin(longitude)]
```

```
dd_train_engineering[, z := sin(latitude)]
```

```
dd_train_engineering[, living_area := supply_area - exclusive_use_area]
```

```
dd_train_engineering[, bathroom_per_living_area := bathroom_count/living_area]
```

```
dd_train_engineering[, area_ratio := exclusive_use_area / supply_area]
```

```
dd_train_engineering[, household_ratio := total_household_count_of_area_type / total_household_count_in_sites]
```

```
dd_train_engineering[, total_household_per_building_count := total_household_count_in_sites / apartment_building_count_in_sites]
```

```
dd_train_engineering[, age_of_apartment := 2021 - year_of_completion]
```

```
dd_train_engineering[, year := as.numeric(substr(transaction_year_month, 1,4))]
```

```
dd_train_engineering[, month := as.numeric(substr(transaction_year_month, 5,6))]
```

```
# feature engineering for test
```

```
dd_test_engineering[, x := cos(latitude) * cos(longitude)]
```

```
dd_test_engineering[, y := cos(latitude) * sin(longitude)]
```

```
dd_test_engineering[, z := sin(latitude)]
```

```
dd_test_engineering[, living_area := supply_area - exclusive_use_area]
```

```
dd_test_engineering[, bathroom_per_living_area := bathroom_count/living_area]
```

```
dd_test_engineering[, area_ratio := exclusive_use_area / supply_area]
```

```
dd_test_engineering[, household_ratio := total_household_count_of_area_type / total_household_count_in_sites]
```

```
dd_test_engineering[, total_household_per_building_count := total_household_count_in_sites / apartment_building_count_in_sites]
```

```
dd_test_engineering[, age_of_apartment := 2021 - year_of_completion]
```

```
dd_test_engineering[, year := as.numeric(substr(transaction_year_month, 1,4))]
```

```
dd_test_engineering[, month := as.numeric(substr(transaction_year_month, 5,6))]
```

Boosting: cross validation without feature engineering

```
boost.price.cv <- gbm(transaction_real_price ~ ., data = train_sample, distribution = "gaussian",
, n.trees = 5000, interaction.depth = 4, cv.folds=10)
```

```
# predict on train dataset
yhat.boost.cv.train = predict(boost.price.cv, newdata = dd_train, n.trees = which.min(boost.price.cv$cv.error))
cat("Train RMSE: ", sqrt(mean((dd_train$transaction_real_price - yhat.boost.cv.train)^2)))
```

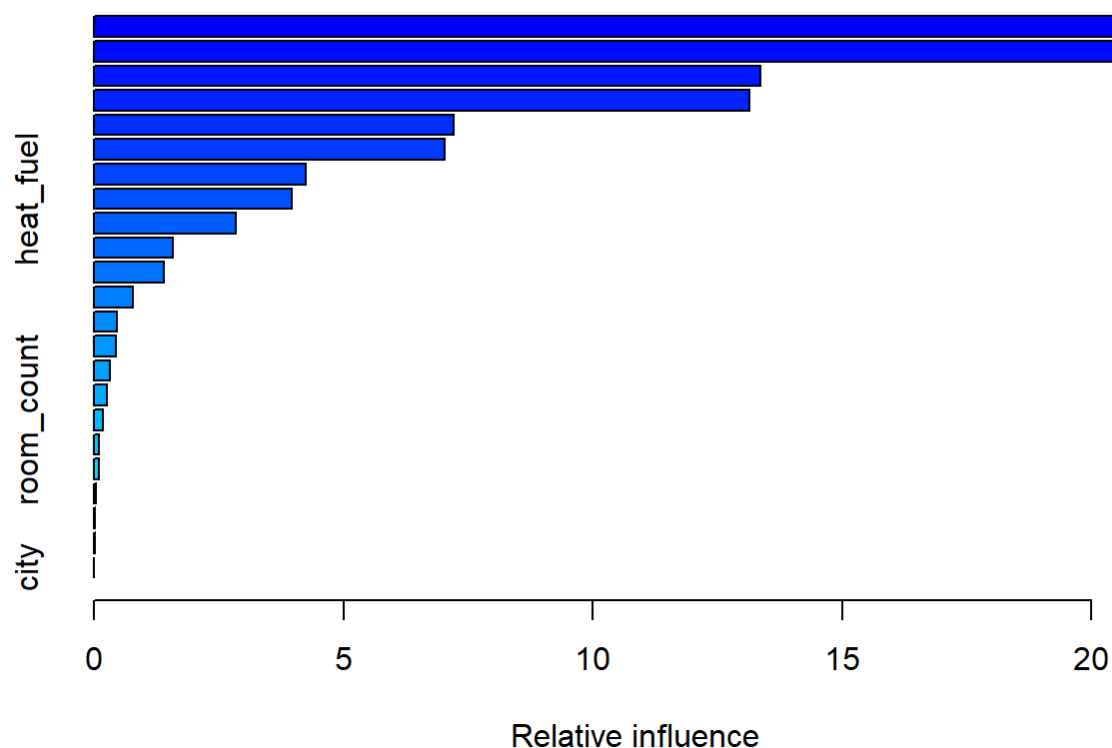
```
## Train RMSE: 44130971
```

```
# predict on the test dataset using the recommended number of trees
predictions <- predict(boost.price.cv, newdata=dd_train, n.trees=which.min(boost.price.cv$cv.error))
cat(" Test RMSE: ", sqrt(mean(dd_train$transaction_real_price - predictions)**2))
```

```
## Test RMSE: 45081.09
```

The most significant features according to the model.

```
summary(boost.price.cv)
```



```
##                                var
## supply_area                   supply_area
## address_by_law                address_by_law
## exclusive_use_area            exclusive_use_area
## latitude                      latitude
## longitude                    longitude
## transaction_year_month        transaction_year_month
## year_of_completion            year_of_completion
## heat_fuel                     heat_fuel
## tallest_building_in_sites     tallest_building_in_sites
## apartment_building_count_in_sites apartment_building_count_in_sites
## total_parking_capacity_in_site total_parking_capacity_in_site
## lowest_building_in_sites      lowest_building_in_sites
## total_household_count_in_sites total_household_count_in_sites
## floor                         floor
## heat_type_district            heat_type_district
## total_household_count_of_area_type total_household_count_of_area_type
## room_count                   room_count
## front_door_structure_corridor front_door_structure_corridor
## bathroom_count               bathroom_count
## transaction_date              transaction_date
## front_door_structure_mixed    front_door_structure_mixed
## heat_type_central             heat_type_central
## city                         city
##                                rel.inf
## supply_area                   21.932027794
## address_by_law                20.524205923
## exclusive_use_area            13.364367188
## latitude                      13.146690161
## longitude                    7.207243203
## transaction_year_month        7.040710045
## year_of_completion            4.247379210
## heat_fuel                     3.966227518
## tallest_building_in_sites     2.850453637
## apartment_building_count_in_sites 1.583823587
## total_parking_capacity_in_site 1.403894078
## lowest_building_in_sites      0.786461733
## total_household_count_in_sites 0.457996352
## floor                         0.434545434
## heat_type_district            0.313210265
## total_household_count_of_area_type 0.267675320
## room_count                   0.176855572
## front_door_structure_corridor 0.103106493
## bathroom_count               0.102814494
## transaction_date              0.033487882
## front_door_structure_mixed    0.027144779
## heat_type_central             0.026619081
## city                         0.003060251
```

select a random sample of 10000 observations

```
train_sample_index = sample(1:nrow(dd_train_engineering),nrow(dd_train_engineering)*4/5)
train_sample = dd_train_engineering[train_sample_index,]
tree.testy_sample <- dd_test_engineering[train_sample_index, transaction_real_price]
```

Boosting: cross validation with feature engineering

```
boost.price.cv <- gbm(transaction_real_price ~ ., data = train_sample, distribution = "gaussian"
, n.trees = 5000, interaction.depth = 4, cv.folds=10)
```

```
yhat.boost.train = predict(boost.price.cv, newdata = dd_train_engineering, n.trees = which.min(b
oost.price.cv$cv.error))
cat("RMSE train: ", sqrt(mean((dd_train_engineering$transaction_real_price - yhat.boost.train)^2
)))
```

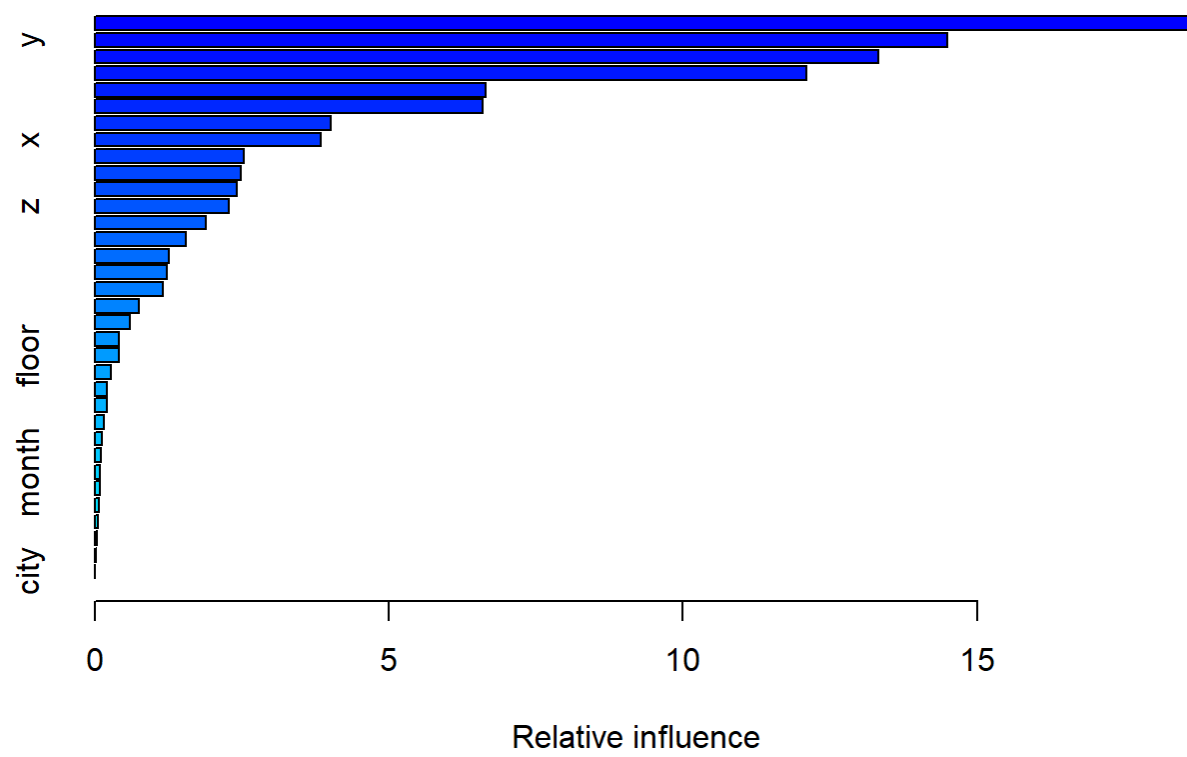
```
## RMSE train: 42532081
```

```
# predict on the test dataset using the recommended number of trees
predictions <- predict(boost.price.cv, newdata = dd_test_engineering, n.trees = which.min(boost.
price.cv$cv.error))
cat(" RMSE test: ", sqrt(mean((dd_test_engineering$transaction_real_price - predictions)**2)))
```

```
## RMSE test: 49999316
```

The most significant features according to the model.

```
summary(boost.price.cv)
```



```

##                                     var
## supply_area                       supply_area
## y                                 y
## exclusive_use_area                exclusive_use_area
## address_by_law                    address_by_law
## latitude                          latitude
## transaction_year_month            transaction_year_month
## heat_fuel                         heat_fuel
## x                                 x
## longitude                          longitude
## year_of_completion                year_of_completion
## tallest_building_in_sites          tallest_building_in_sites
## z                                 z
## living_area                       living_area
## total_household_per_building_count total_household_per_building_count
## apartment_building_count_in_sites apartment_building_count_in_sites
## total_parking_capacity_in_site     total_parking_capacity_in_site
## age_of_apartment                  age_of_apartment
## heat_type_district                heat_type_district
## lowest_building_in_sites           lowest_building_in_sites
## area_ratio                        area_ratio
## floor                             floor
## total_household_count_in_sites     total_household_count_in_sites
## year                              year
## bathroom_per_living_area           bathroom_per_living_area
## household_ratio                    household_ratio
## total_household_count_of_area_type total_household_count_of_area_type
## room_count                         room_count
## month                             month
## bathroom_count                     bathroom_count
## front_door_structure_corridor       front_door_structure_corridor
## heat_type_central                   heat_type_central
## transaction_date                    transaction_date
## front_door_structure_mixed          front_door_structure_mixed
## city                               city
##                                   rel.inf
## supply_area                       18.608857749
## y                                 14.494912540
## exclusive_use_area                13.326192767
## address_by_law                    12.105804529
## latitude                          6.648128023
## transaction_year_month            6.596668873
## heat_fuel                         4.006117385
## x                                 3.835406950
## longitude                          2.538490247
## year_of_completion                2.491341887
## tallest_building_in_sites          2.413669663
## z                                 2.272997050
## living_area                       1.888397151
## total_household_per_building_count 1.556343773
## apartment_building_count_in_sites 1.260717839
## total_parking_capacity_in_site     1.227589358
## age_of_apartment                  1.158348716

```

```
## heat_type_district      0.758079011
## lowest_building_in_sites 0.590870920
## area_ratio              0.406922451
## floor                   0.405400016
## total_household_count_in_sites 0.269270636
## year                    0.212812289
## bathroom_per_living_area 0.202415716
## household_ratio          0.154866778
## total_household_count_of_area_type 0.122390273
## room_count              0.098498882
## month                   0.085613219
## bathroom_count          0.083490173
## front_door_structure_corridor 0.072583439
## heat_type_central        0.049241930
## transaction_date         0.034340784
## front_door_structure_mixed 0.021255344
## city                     0.001963639
```